

Distributed Systems

Subheadline

► Stefan Henkler

E-Mail: stefan.henkler@hshl.de

► OpenMP

- Explicit thread management is too intrusive
- Explicit thread management is often unnecessary
 - Consider the matrix multiplication
- Solution
 - OpenMP is a nonintrusive framework for parallel programming

► What is OpenMP

- Open specifications for multiprocessing via collaborative work between interested parties from the hardware and software industry, government and academia.
- OpenMP is an Application Program Interface (API) that may be used to explicitly direct ***multi-threaded, shared memory parallelism***.
 - API components: Compiler Directives, Runtime Library Routines. Environment Variables
- OpenMP is a directive-based method to invoke parallel computations on share-memory multiprocessors

► What is OpenMP?

- OpenMP API is specified for C/C++ and Fortran.
- OpenMP is not intrusive to the original serial code: instructions appear in comment statements for fortran and pragmas for C/C++.
- OpenMP website: <http://www.openmp.org>
 - Materials in this lecture are taken from various OpenMP tutorials on the website and other places.

► Why OpenMP?

- OpenMP is portable: supported by HP, IBM, Intel, SGI, SUN, and others
 - It is the de facto standard for writing shared memory programs.
- OpenMP can be implemented incrementally

► How to compile and run OpenMP programs?

► Gcc 4.2 and above supports OpenMP 3.0


► `gcc -fopenmp a.c`

► To run: 'a.out'

► To change the number of threads:

► `setenv OMP_NUM_THREADS 4 (tcsh)` or `export OMP_NUM_THREADS=4(bash)`

OpenMP execution model



- ▶ OpenMP uses the fork-join model of parallel execution.
 - ▶ All OpenMP programs begin with a single **master thread**.
 - ▶ The master thread executes sequentially until a **parallel region** is encountered, when it creates a **team of parallel threads** (FORK).
 - ▶ When the team threads complete the parallel region, they synchronize and terminate, leaving only the master thread that executes sequentially (JOIN).

► OpenMP general code structure

```
#include <omp.h>
```

```
main () {
```

```
    int var1, var2, var3;
```

```
    Serial code
```

```
    . . .
```

```
    /* Beginning of parallel section. Fork a team of threads. Specify variable scoping*/
```

```
    #pragma omp parallel private(var1, var2) shared(var3)
```

```
    {
```

```
        /* Parallel section executed by all threads */
```

```
        . . .
```

```
        /* All threads join master thread and disband*/
```

```
    }
```


```
    Resume serial code
```

```
    . . .
```

```
}
```


Data model

- Private and shared variables



P = private data space
G = global data space

- Variables in the global data space are accessed by all parallel threads (**shared** variables).
- Variables in a thread's private space can only be accessed by the thread (**private** variables)

► OpenMP directives

► Format:

`#pragma omp directive-name [clause,...] newline`
(use `'\'` for multiple lines)

► Example:

`#pragma omp parallel default(shared) private(beta,pi)`

► Scope of a directive is one block of statements { ...}

► Parallel region construct

- A block of code that will be executed by multiple threads.

```
#pragma omp parallel [clause ...]  
{  
    .....  
} (implied barrier)
```


Clauses: *if (expression)*, *private (list)*, *shared (list)*, *default (shared | none)*,
reduction (operator: list), *firstprivate(list)*, *lastprivate(list)*

- *if (expression)*: only in parallel if expression evaluates to true
- *private(list)*: everything private and local (no relation with variables outside the block).
- *shared(list)*: data accessed by all threads
- *default (none|shared)*

Loop - index

```
#pragma omp parallel for private( privIndx, privDbl )  
for ( i = 0; i < arraySize; i++ ) {  
    for ( privIndx = 0; privIndx < 16; privIndx++ ) {  
        privDbl = ( (double) privIndx ) / 16;  
        y[i] = sin( exp( cos( - exp( sin(x[i]) ) ) ) ) ) + cos(  
            privDbl );  
    }  
}
```

Parallel for loop index is
Private by default.



execution context for "arrayUpdate_II"

► Reduction

```
Sum = 0.0;
#pragma parallel default(none) shared (n, x) private (l) reduction(+ : sum)
{
    For(l=0; l<n; l++) sum = sum + x(l);
}
```

- Updating sum must avoid racing condition
- With the reduction clause, OpenMP generates code such that the race condition is avoided.

▶ Work-sharing constructs

- ▶ `#pragma omp for [clause ...]`
 - ▶ ... as discussed
- ▶ `#pragma omp section [clause ...]`
 - ▶ Only one thread per section
- ▶ `#pragma omp single [clause ...]`
 - ▶ Only one thread executes block.
 - ▶ If team of threads: All other ones wait until single block is executed
- ▶ The work is distributed over the threads
- ▶ Must be enclosed in parallel region

The omp for directive: example

Disable synchronization after for

```
#pragma omp parallel default(none) \
    shared(n,a,b,c,d) private(i)
{
    #pragma omp for nowait
    for (i=0; i<n-1; i++)
        b[i] = (a[i] + a[i+1])/2;

    #pragma omp for nowait
    for (i=0; i<n; i++)
        d[i] = 1.0/c[i];

} /*-- End of parallel region --*/
    (implied barrier)
```

The omp sections clause - example

```
#pragma omp parallel default(none) \  
    shared(n,a,b,c,d) private(i)  
{  
    #pragma omp sections nowait  
    {  
        #pragma omp section  
        for (i=0; i<n-1; i++)  
            b[i] = (a[i] + a[i+1])/2;  
  
        #pragma omp section  
        for (i=0; i<n; i++)  
            d[i] = 1.0/c[i];  
  
    } /*-- End of sections --*/  
  
} /*-- End of parallel region --*/
```


► Synchronization: barrier

```
For(l=0; l<N; l++)  
    a[l] = b[l] + c[l];
```

```
For(l=0; l<N; l++)  
    d[l] = a[l] + b[l]
```

Both loops are in parallel region
With no synchronization in between.
Problem: dependencies between both loops
-> race condition!

Add explicit barrier!

```
For(l=0; l<N; l++)  
    a[l] = b[l] + c[l];  
  
#pragma omp barrier  
  
For(l=0; l<N; l++)  
    d[l] = a[l] + b[l]
```

► Critical section

```
For(I=0; I<N; I++) {  
    .....  
    sum += A[I];  
    .....  
}
```

Cannot be parallelized if sum is shared.

Fix:

```
For(I=0; I<N; I++) {  
    .....  
    #pragma omp critical  
    {  
        sum += A[I];  
    }  
    .....  
}
```

- ▶ **OpenMP environment variables**

- ▶ OMP_NUM_THREADS
- ▶ OMP_SCHEDULE

► OpenMP runtime environment

- `omp_get_num_threads`
- `omp_get_thread_num`
- `omp_in_parallel`
- ...

► Sequential Matrix Multiply

```
For (l=0; l<n; l++)  
    for (j=0; j<n; j++)  
        c[l][j] = 0;  
        for (k=0; k<n; k++)  
            c[l][j] = c[l][j] + a[l][k] * b[k][j];
```

► OpenMP Matrix Multiply

```
#pragma omp parallel for private(j, k)
For (l=0; l<n; l++)
    for (j=0; j<n; j++)
        c[l][j] = 0;
        for (k=0; k<n; k++)
            c[l][j] = c[l][j] + a[l][k] * b[k][j];
```

► Summary

- OpenMP provides a compact, yet powerful programming model for shared memory programming
 - It is very easy to use OpenMP to create parallel programs.
- OpenMP preserves the sequential version of the program
- Developing an OpenMP program:
 - Start from a sequential program
 - Identify the code segment that takes most of the time.
 - Determine whether the important loops can be parallelized
 - The loops may have critical sections, reduction variables, etc
 - Determine the shared and private variables.
 - Add directives