

Distributed Systems


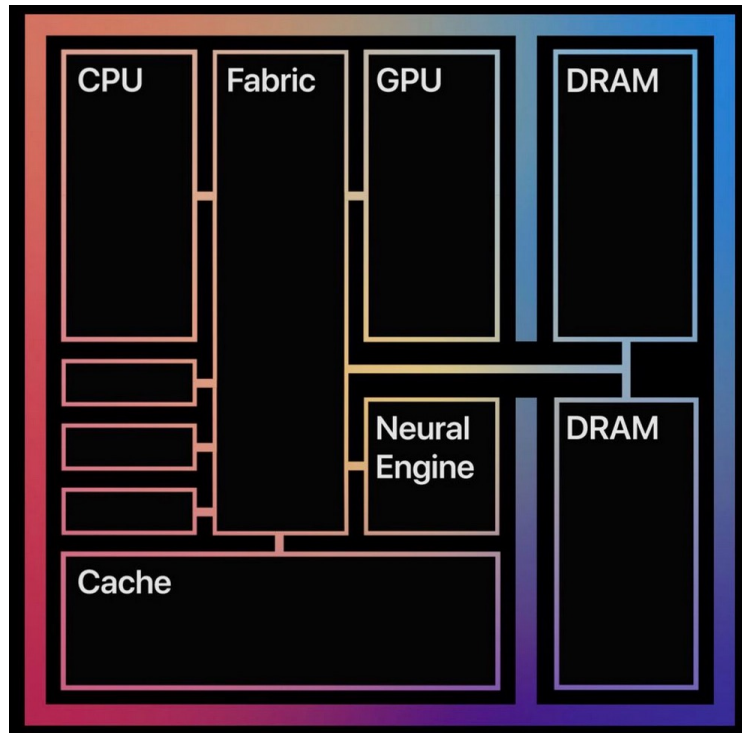
Subheadline

► Stefan Henkler

E-Mail: stefan.henkler@hshl.de



► Motivation

► What is that?



► Motivation

Enabler for



Artificial Intelligence

Million-dollar babies

The Economist

As Silicon Valley fights for talent, universities struggle to hold on to their stars

WIRED

ARTIFICIAL INTELLIGENCE FINALLY ENTERED OUR
EVERYDAY WORLD



Toyota Invests \$1 Billion in Artificial Intelligence in U.S.

Autonomes Fahren in Deutschland

Bosch schickt bald Robo-Taxis auf die Straße



Microsoft and Baidu partner to bring AI powered
vehicles to the automotive market

Bloomberg
Business

The First Person to Hack
the iPhone Built a
Self-Driving Car. In His
Garage




► Overview

- I. **Introduction - parallel architectures**
- II. Distributed Architectures
- III. GPU architecture and programming

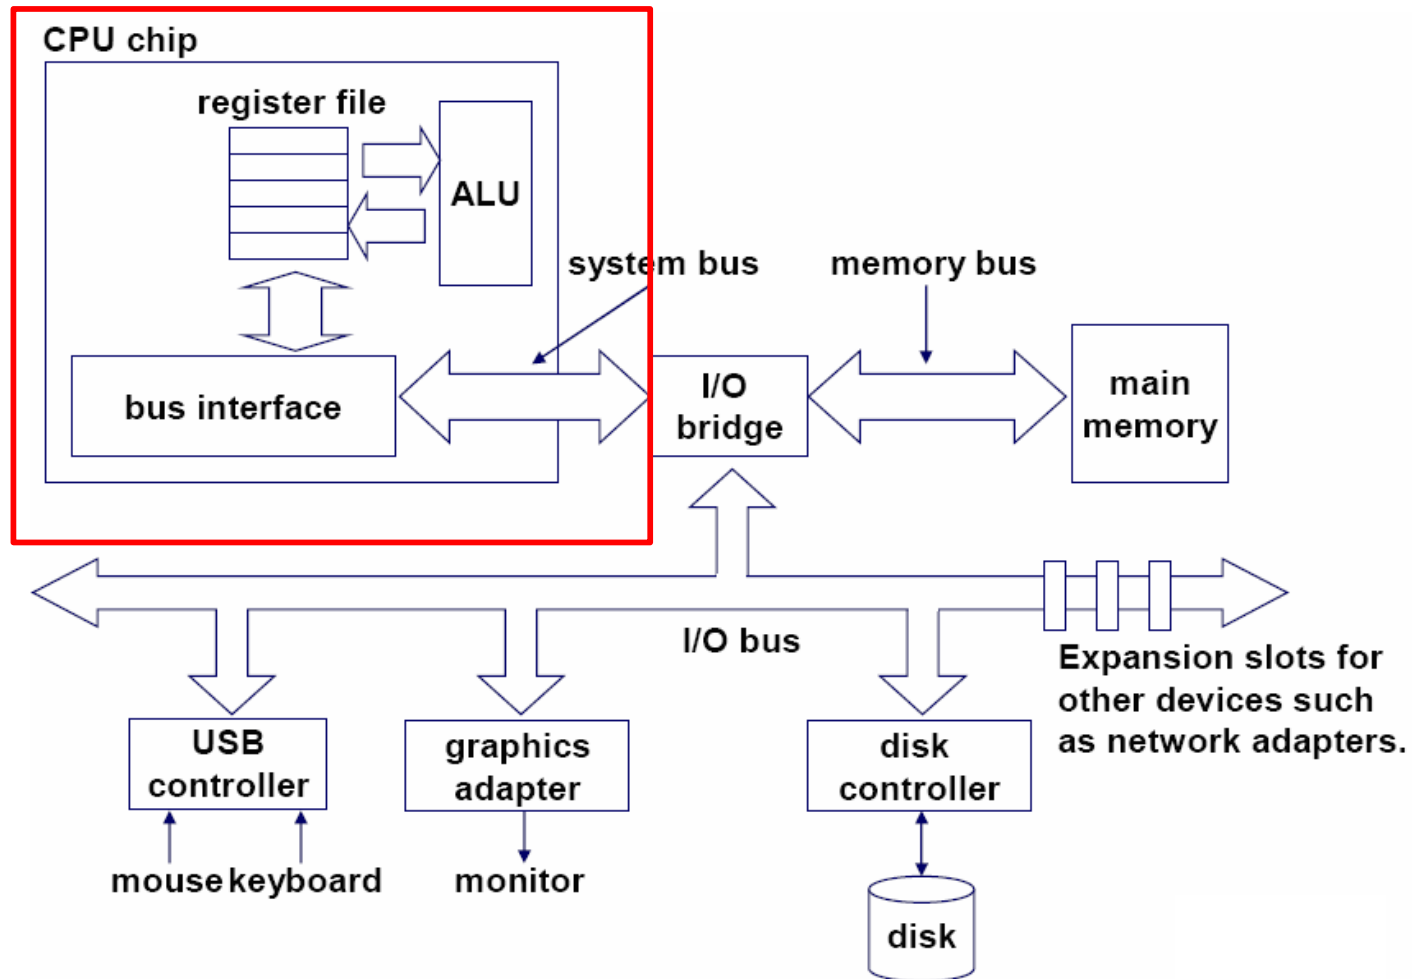
► Multi-core

The secret is under the hood

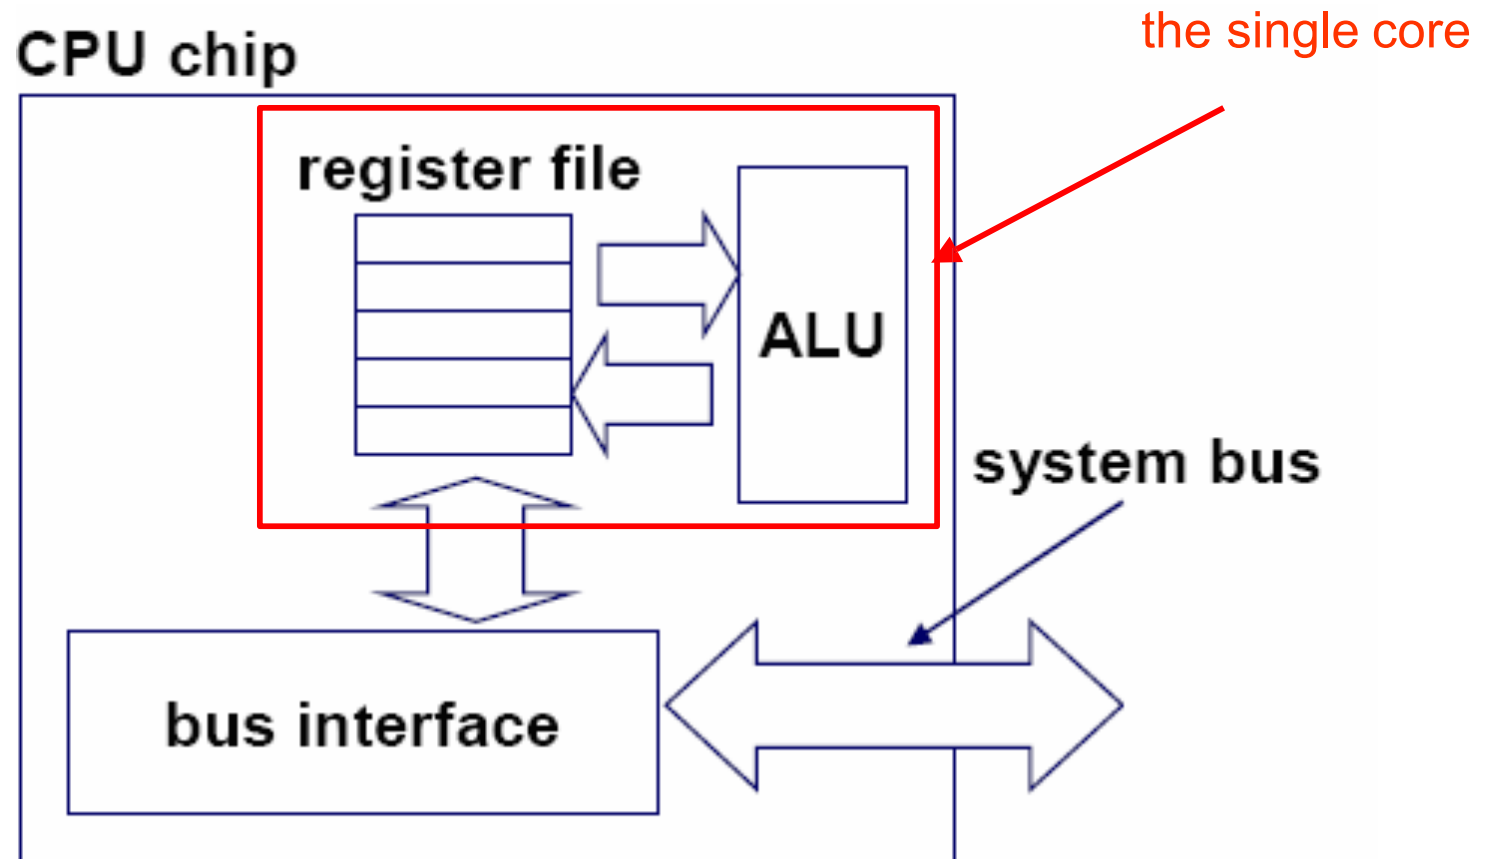


<https://www.amd.com/en/processors/epyc-7002-series>

Single-core computer

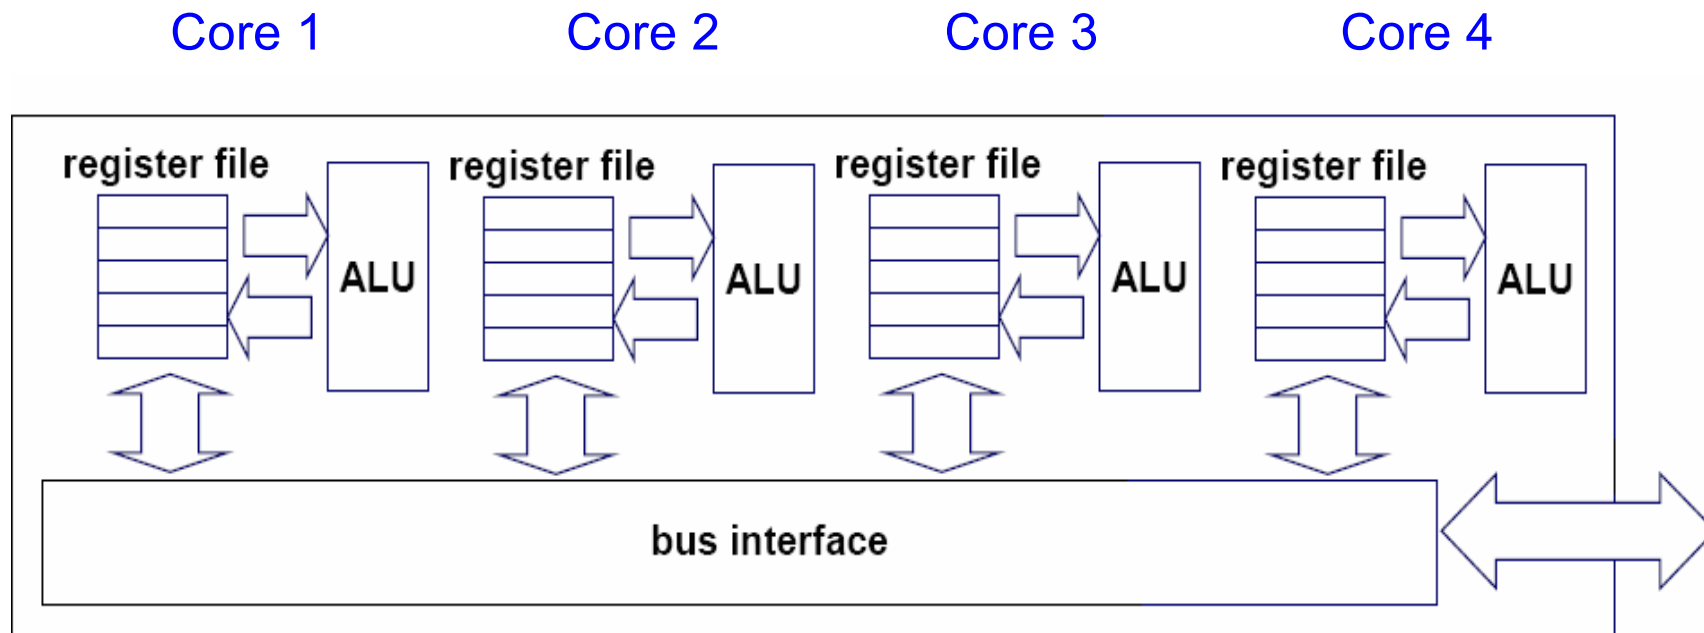


Single-core CPU chip



Multi-core architectures


Replicate multiple processor cores on a single die.



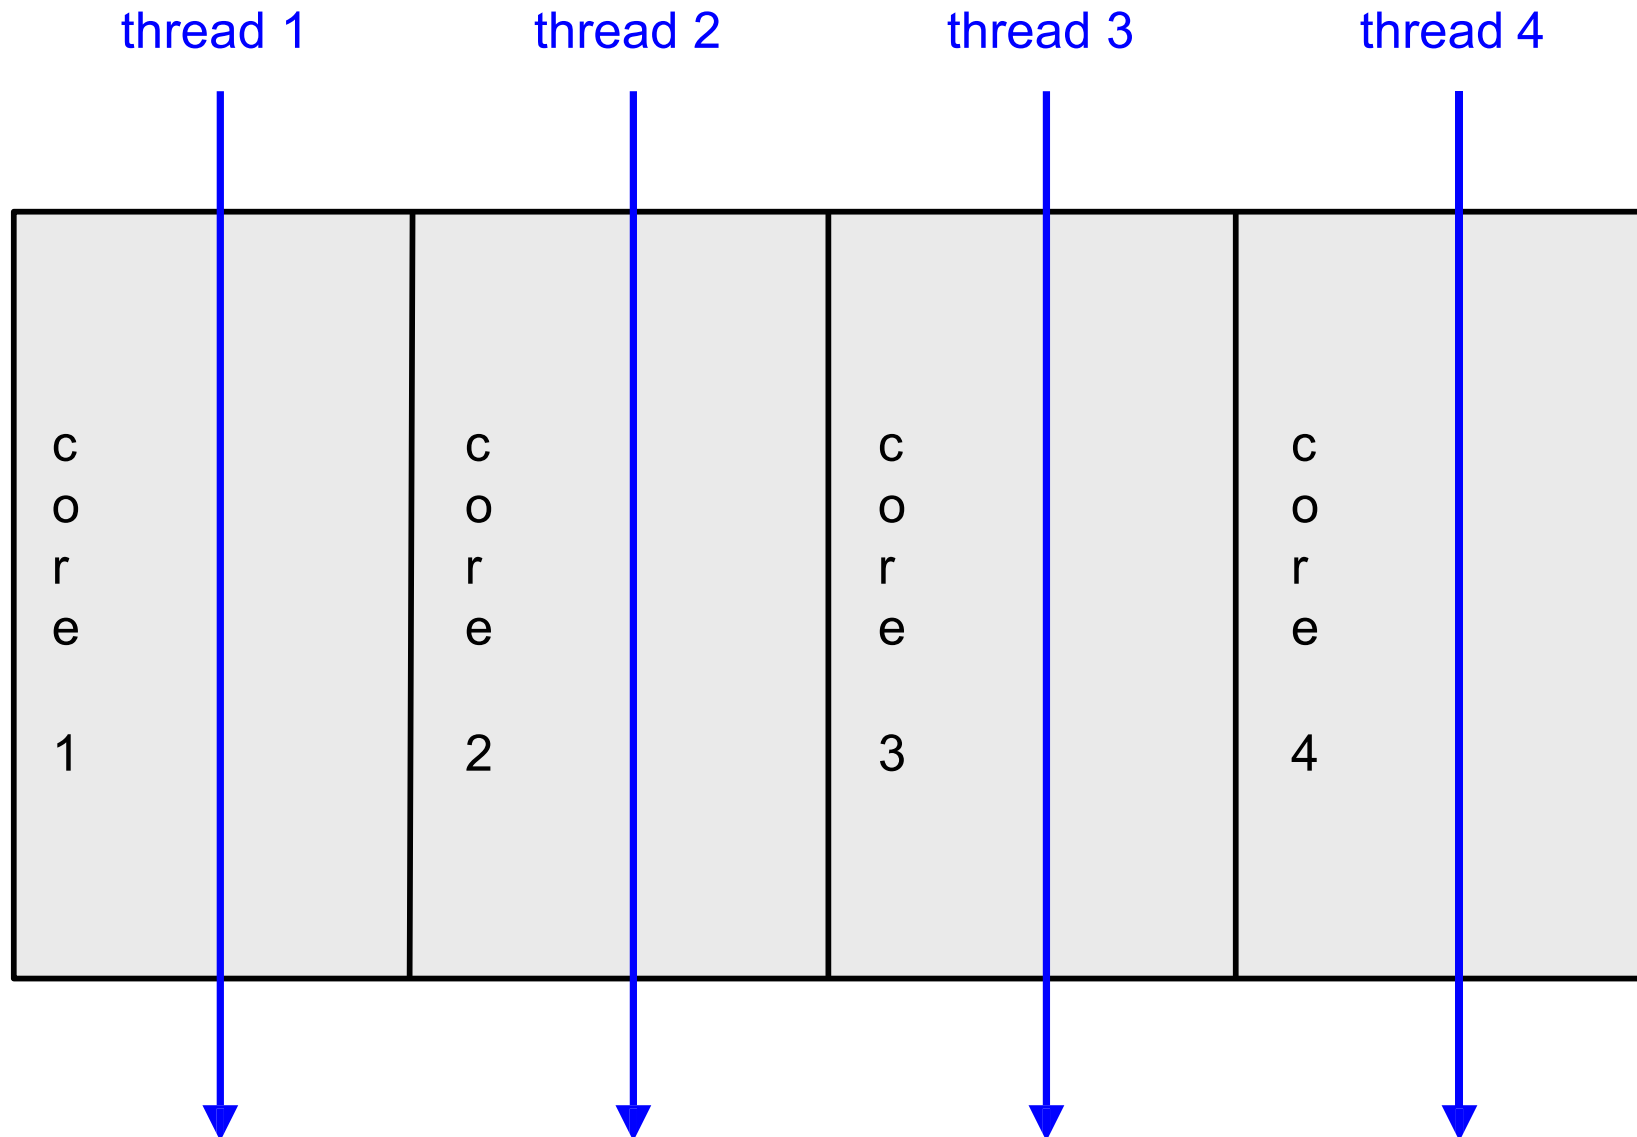
Multi-core CPU chip

Multi-core CPU chip


- The cores fit on a single processor socket
- Also called CMP (Chip Multi-Processor)



The cores run in parallel



Multicore and time slicing



- OS perceives each core as a separate processor
- OS scheduler maps threads/processes to different cores
- Most major OS support multi-core today

▶ Instruction-level parallelism

- Parallelism at the machine-instruction level
- The processor can re-order, pipeline instructions, split them into microinstructions, do aggressive branch prediction, etc.
- Instruction-level parallelism enabled rapid increases in processor speeds over the last 15 years

▶ Thread-level parallelism (TLP)

- This is parallelism on a more coarser scale
- E.g. server can serve each client in a separate thread (Web server, database server)
- A computer game can do AI, graphics, and physics in three separate threads
- Single-core superscalar processors cannot fully exploit TLP
- Multi-core architectures are the next step in processor evolution: explicitly exploiting TLP

General context: Multiprocessors

- Multiprocessor is any computer with several processors
- SIMD
 - Single instruction, multiple data
 - Modern graphics cards
- MIMD
 - Multiple instructions, multiple data



Lemieux cluster,
Pittsburgh
supercomputing
center

▶ Multiprocessor memory types

- Shared memory
 - one (large) common shared memory for all processors
- Distributed memory
 - each processor has its own (small) local memory
 - its content is not replicated anywhere else

Multi-core processor

Special kind of a multiprocessor

- All processors are on the same chip
- Multi-core processors are MIMD: Different cores execute different threads (Multiple Instructions), operating on different parts of memory (Multiple Data).
- Multi-core is a shared memory multiprocessor: All cores share the same memory

Examples

- Database servers
- Web servers (Web commerce)
- Compilers
- Multimedia applications
- Scientific applications, CAD/CAM
- In general, applications with *Thread-level parallelism* (as opposed to instruction- level parallelism)

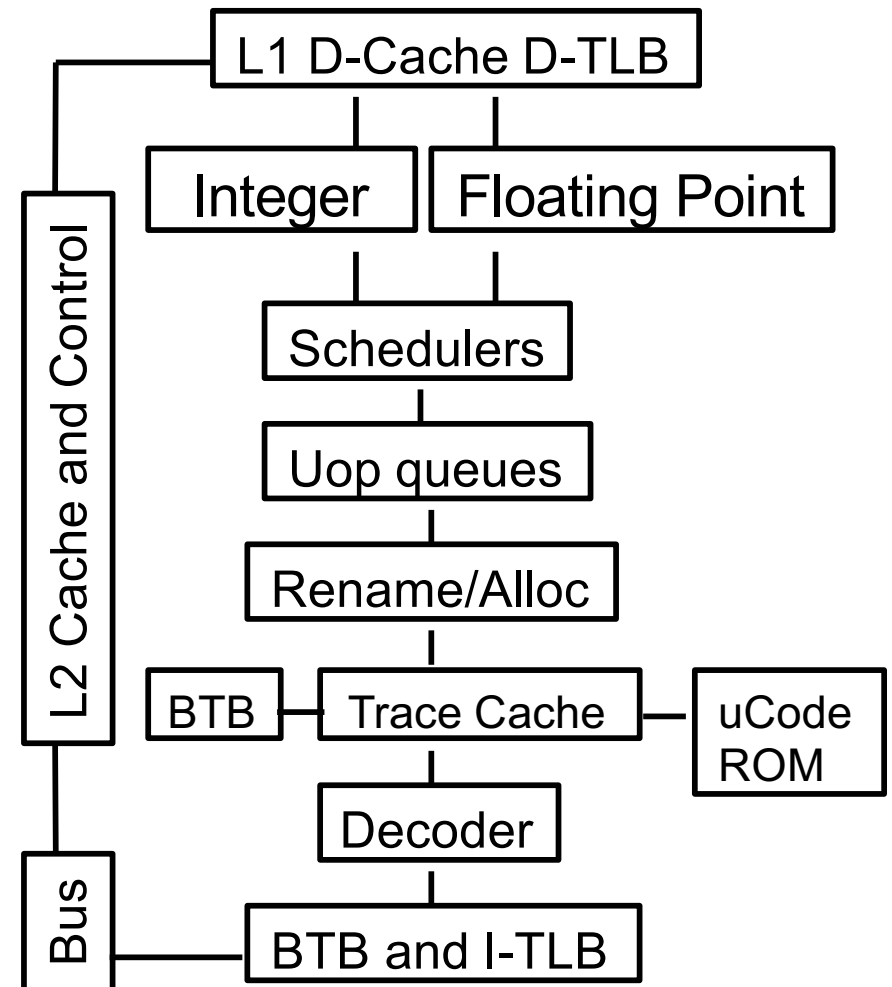
More examples

- Editing a photo while recording a TV show through a digital video recorder
- Downloading software while running an anti-virus program
- “Anything that can be threaded today will map efficiently to multi-core”
- BUT: some applications difficult to parallelize

Simultaneous multithreading

- Problem (example)
 - Waiting for the result of a long floating point (or integer) operation
 - Waiting for data to arrive from memory

Other execution units wait unused

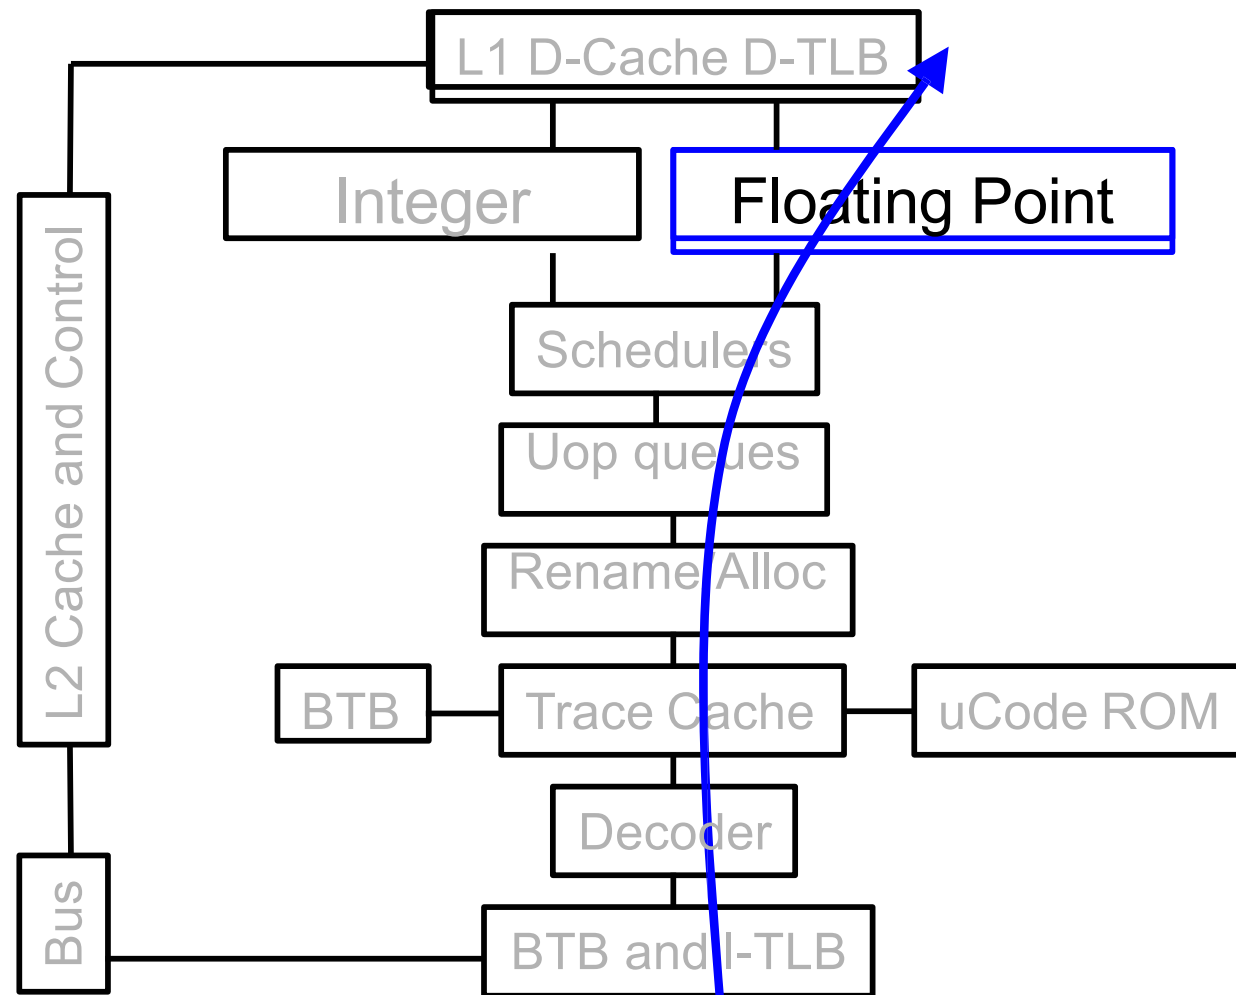


Source: Intel

► Simultaneous multithreading (SMT)

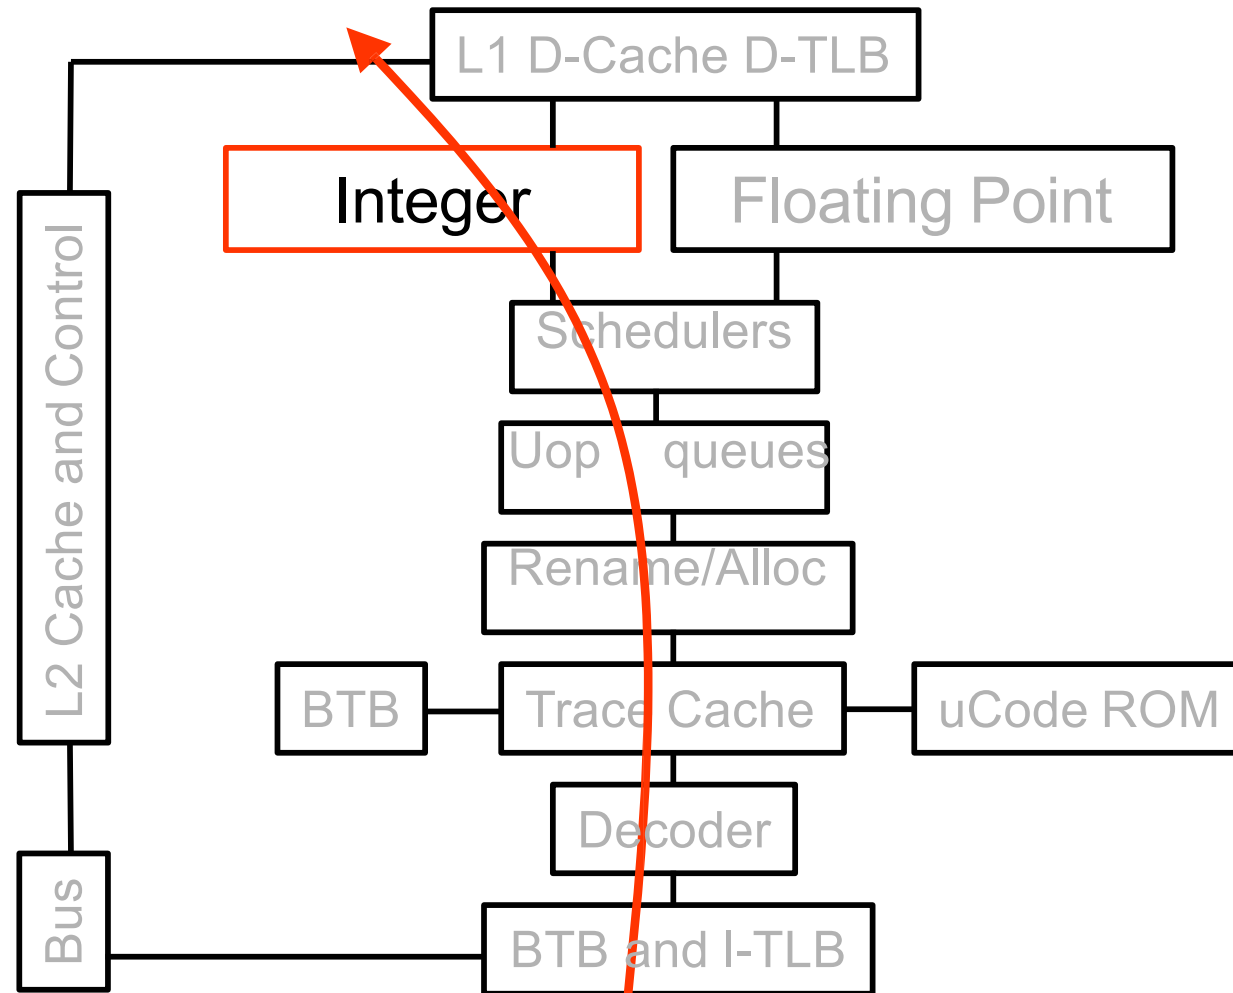
- Permits multiple independent threads to execute **SIMULTANEOUSLY** on the **SAME** core
- Weaving together multiple “threads” on the same core
- Example: if one thread is waiting for a floating point operation to complete, another thread can use the integer units

Without SMT, only a single thread can run at any given time



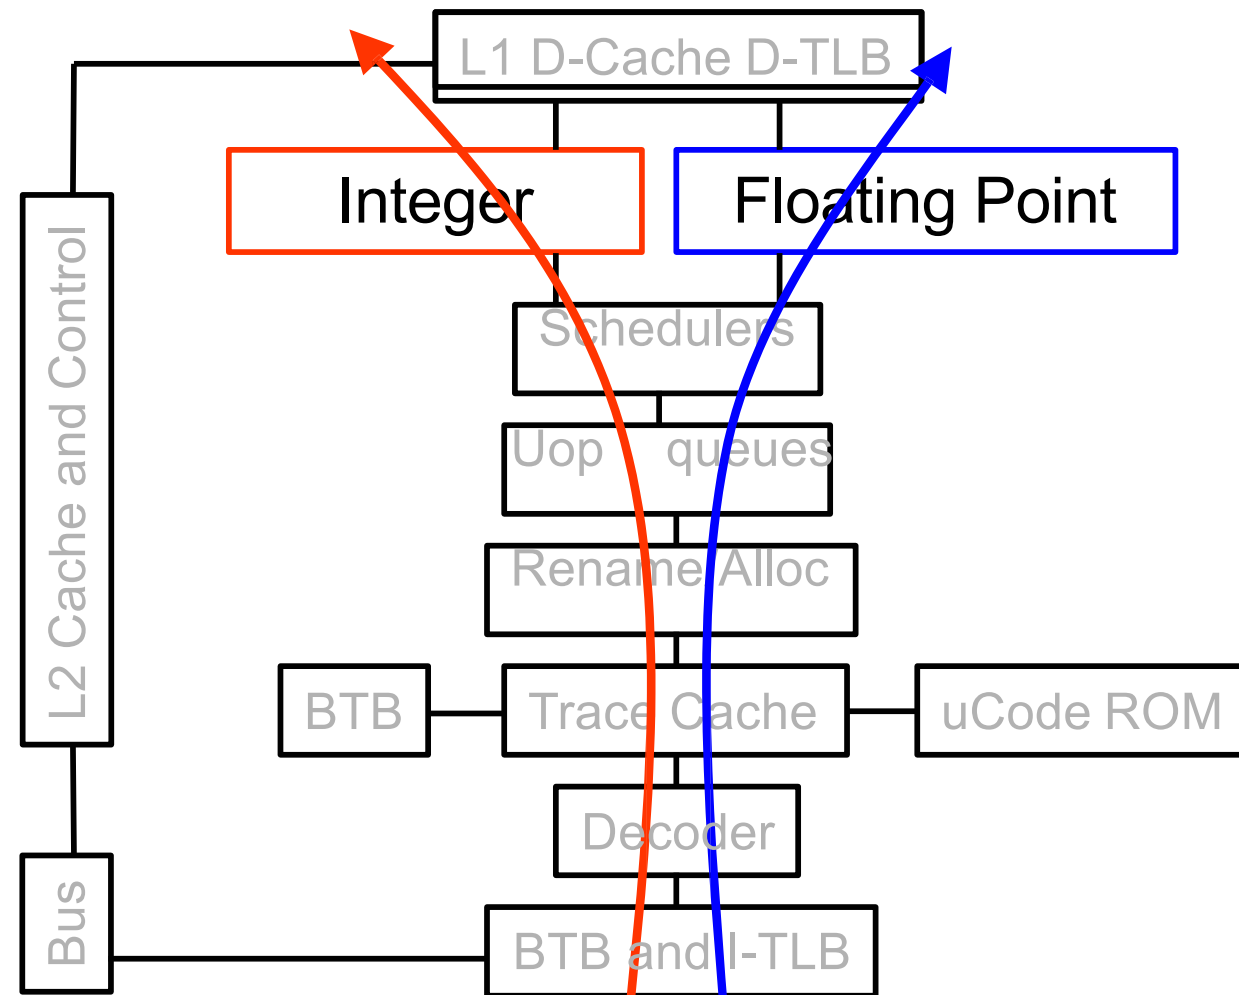
Thread 1: floating point

Without SMT, only a single thread can run at any given time



Thread 2:
integer operation

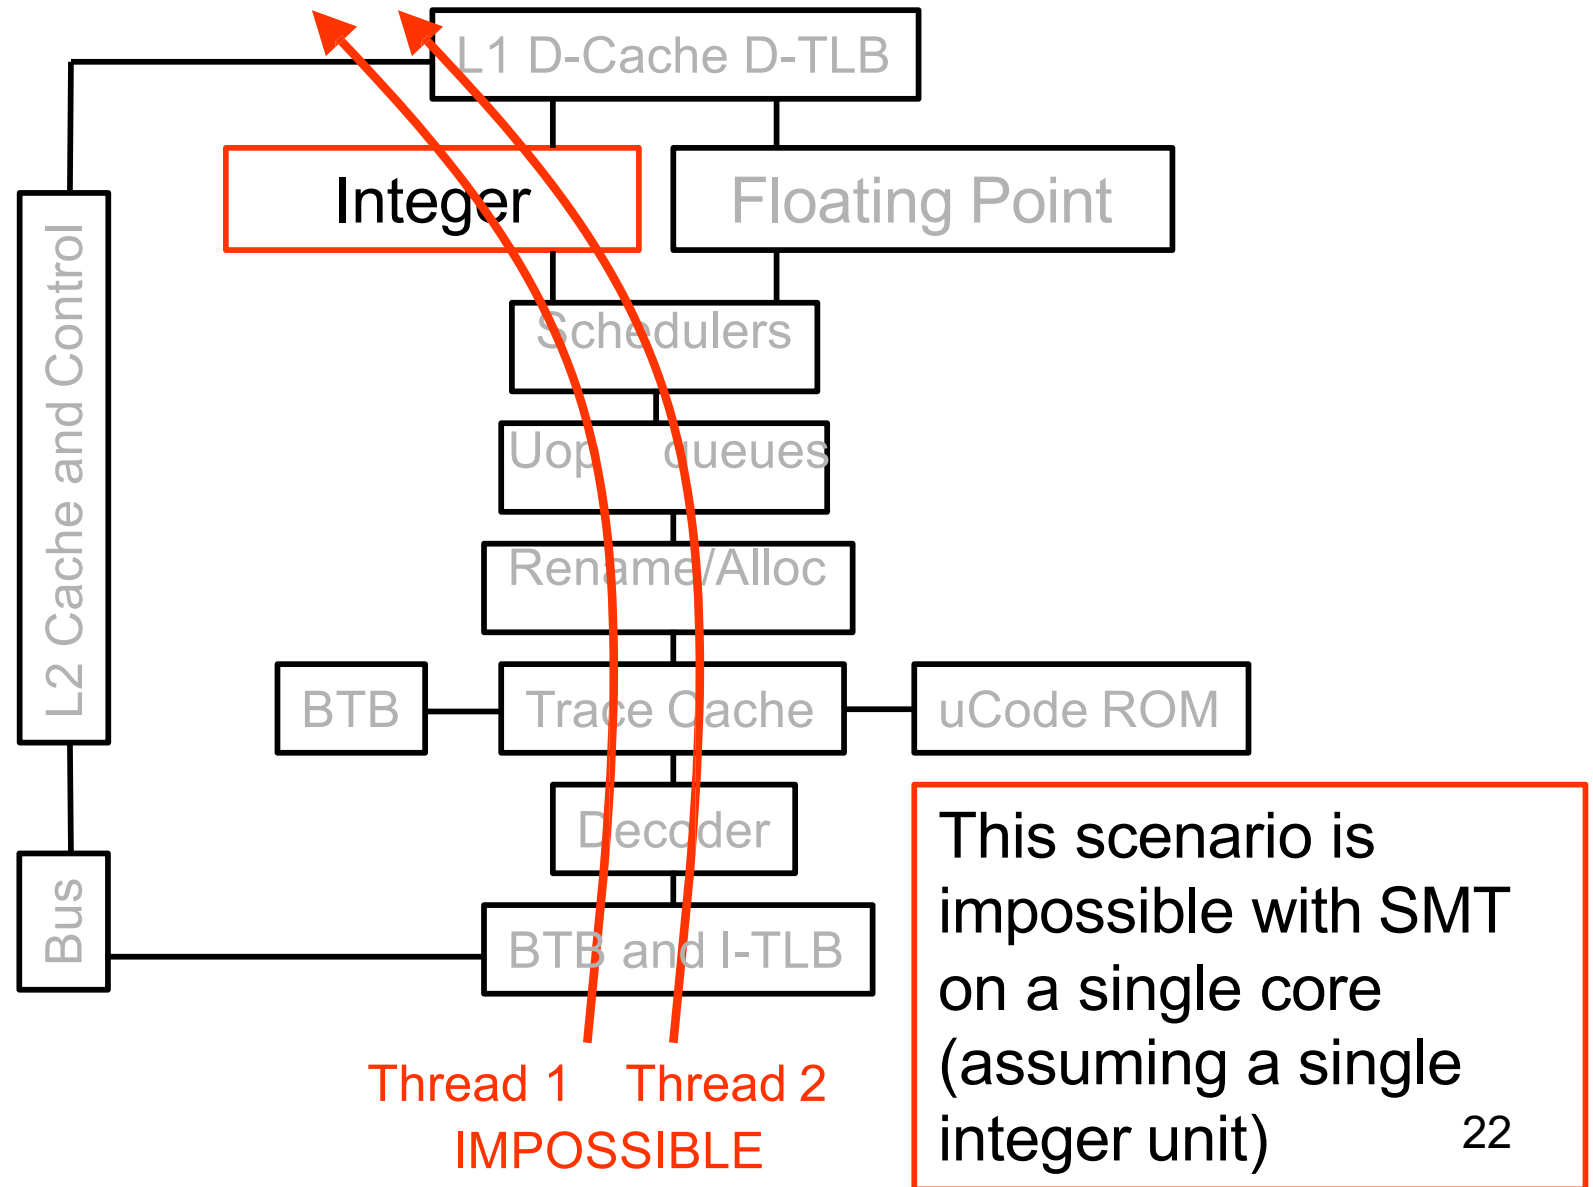
SMT processor: both threads can run concurrently



Thread 2:
integer operation

Thread 1: floating point

But: Can't simultaneously use the same functional unit

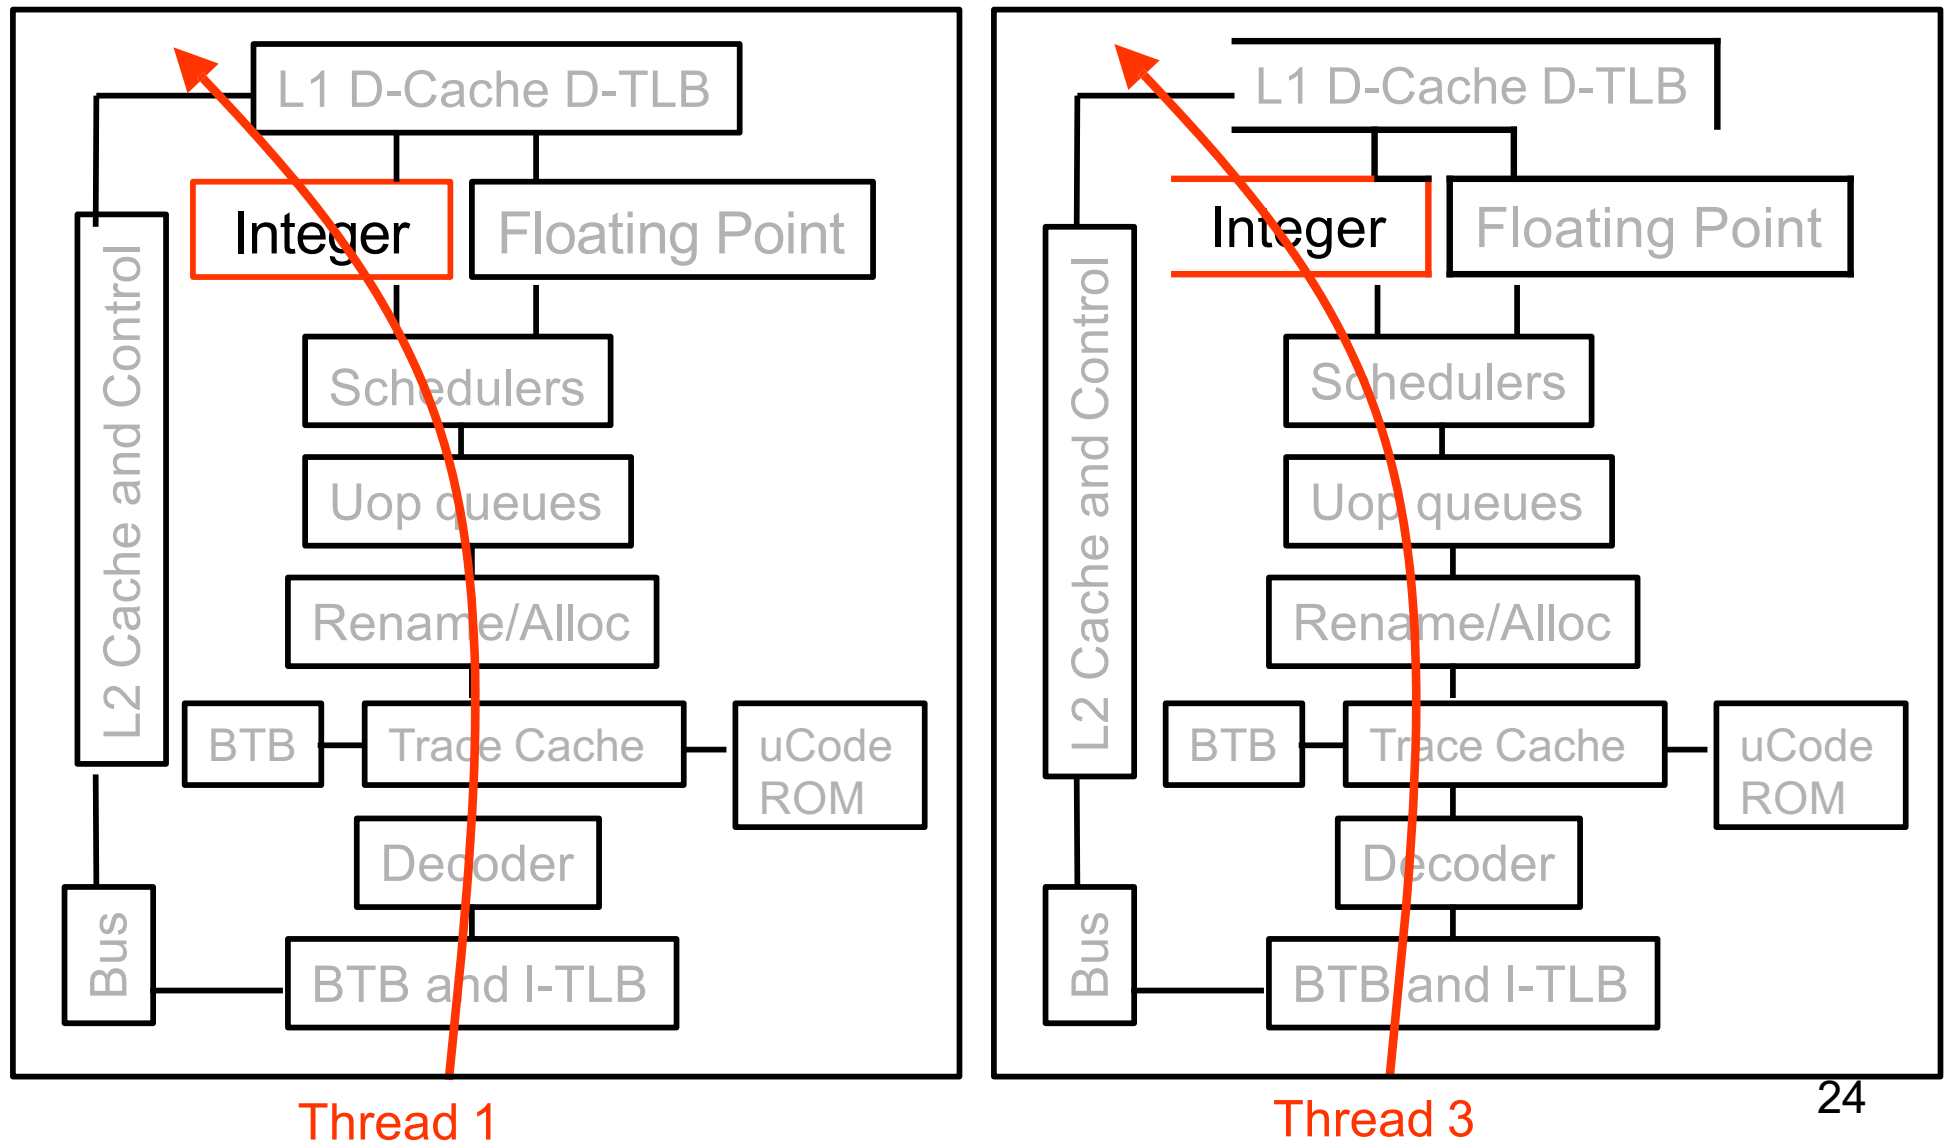


SMT not a “true” parallel processor

- Enables better threading (e.g. up to 30%)
- OS and applications perceive each simultaneous thread as a separate “virtual processor”
- The chip has only a single copy of each resource
- Compare to multi-core:
each core has its own copy of resources

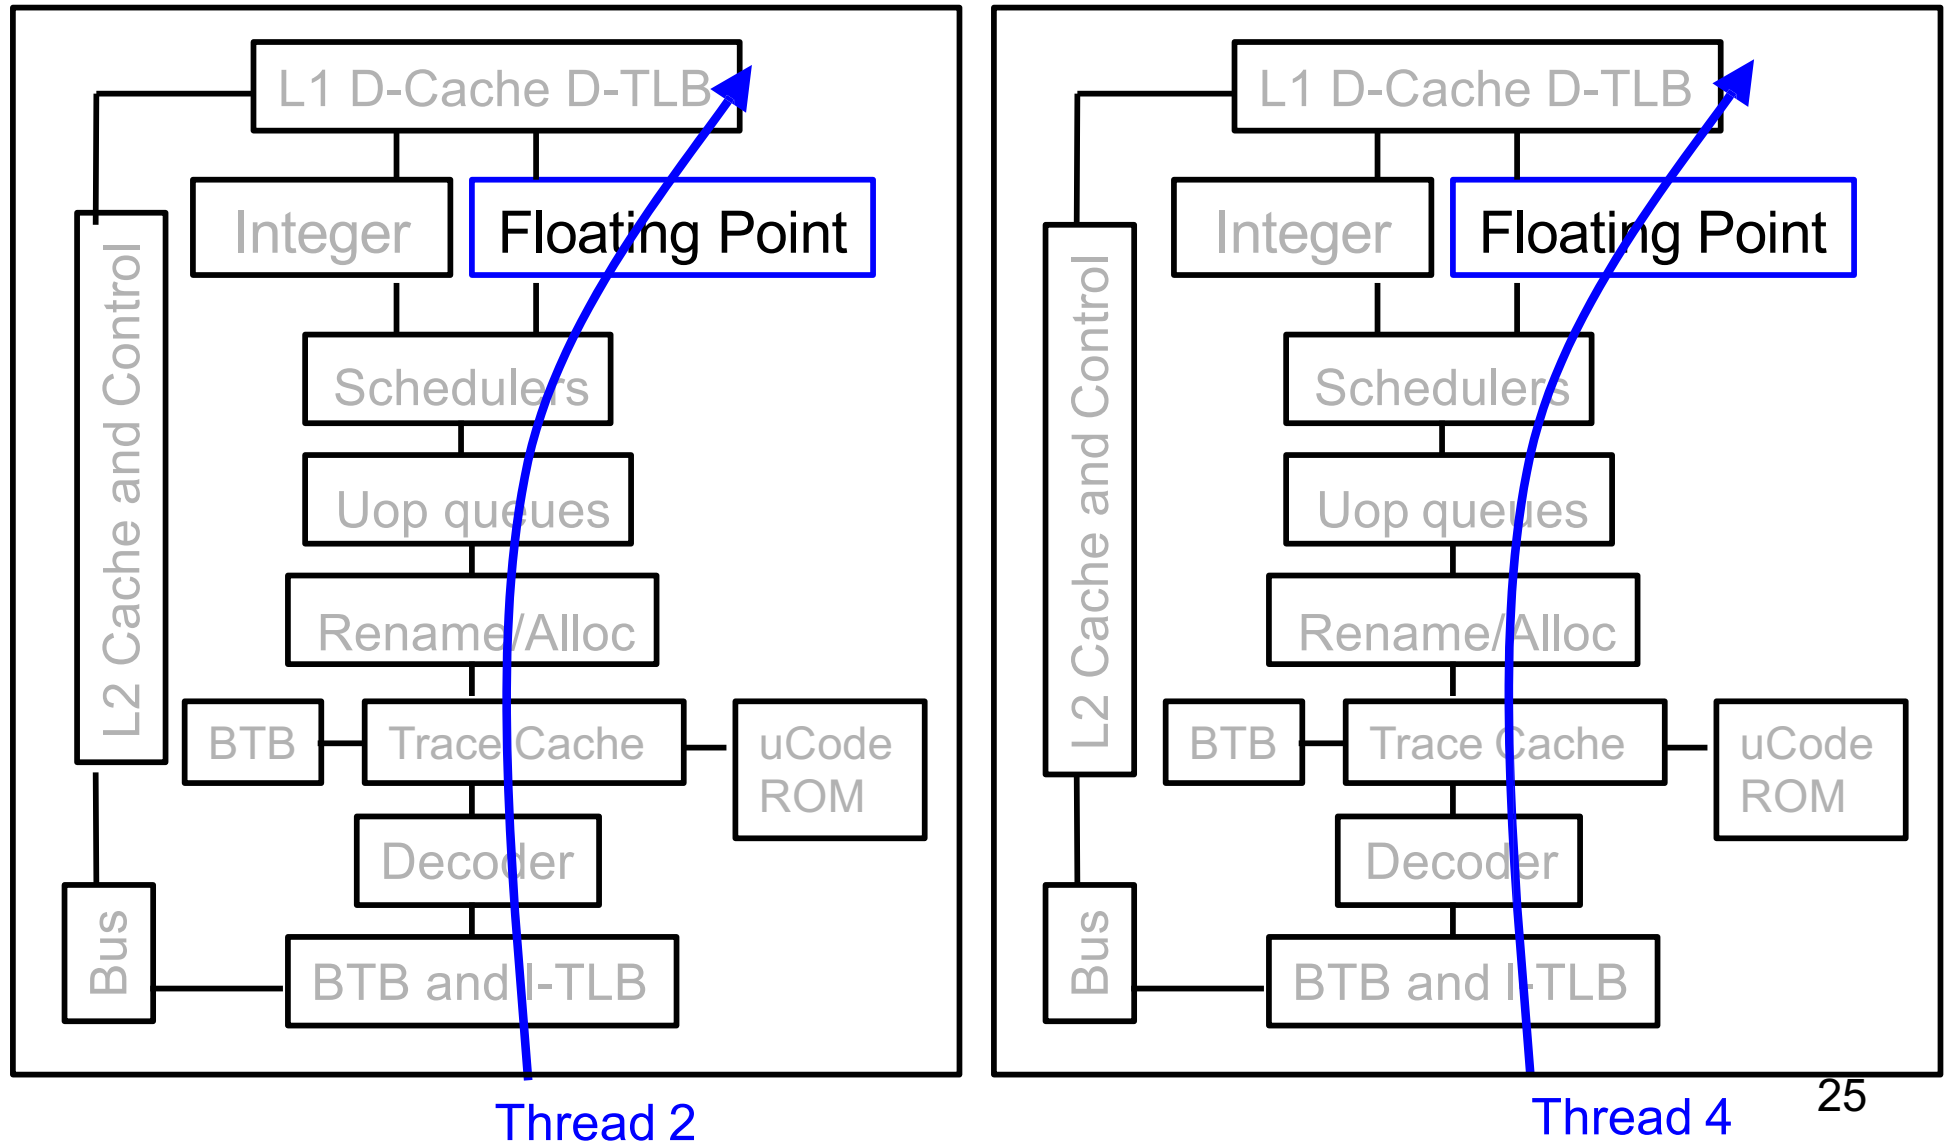
Multi-core:

threads can run on separate cores



Multi-core:

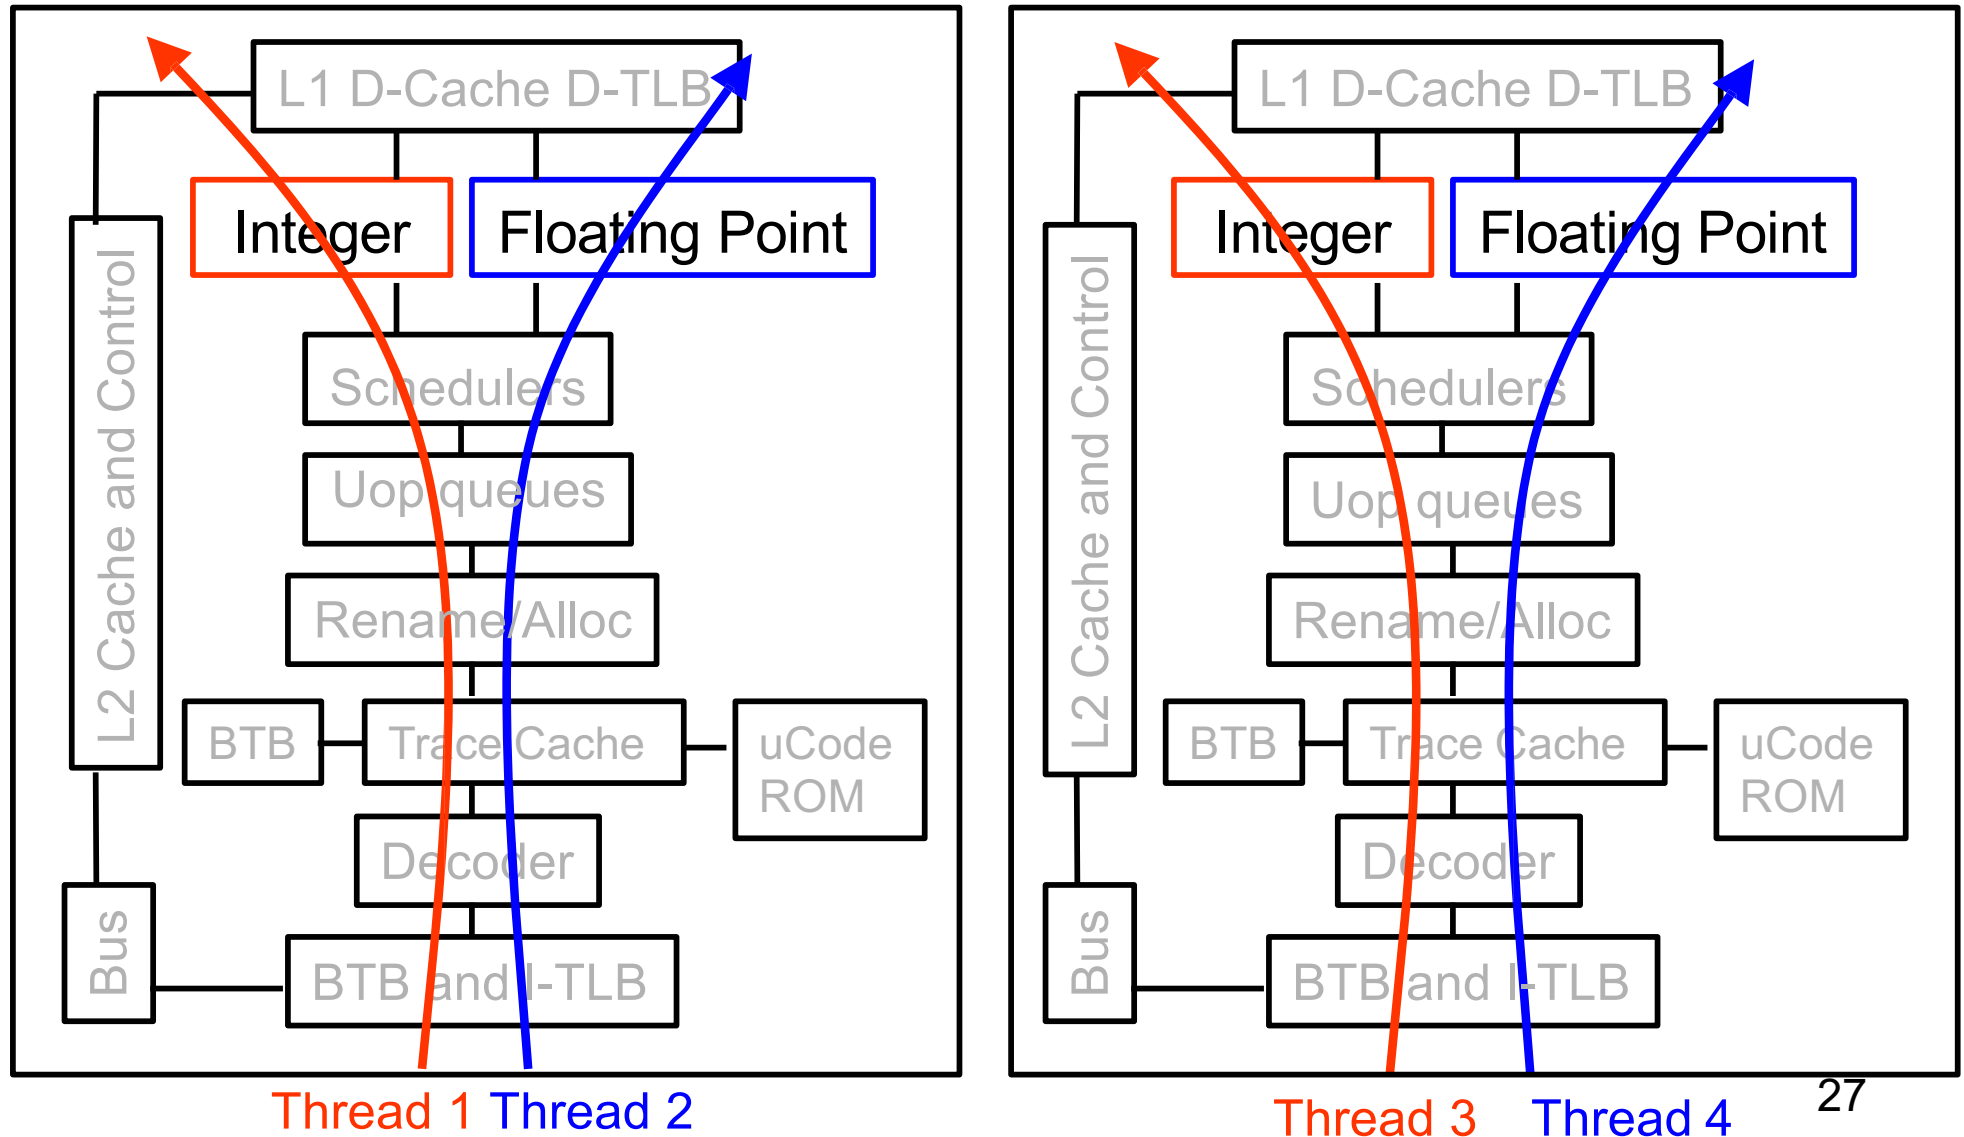
threads can run on separate cores



Combining Multi-core and SMT

- Cores can be SMT-enabled (or not)
- The different combinations:
 - Single-core, non-SMT: standard uniprocessor
 - Single-core, with SMT
 - Multi-core, non-SMT
 - Multi-core, with SMT: our fish machines
- The number of SMT threads:
2, 4, or sometimes 8 simultaneous threads
- Intel calls them “hyper-threads”

SMT Dual-core: all four threads can run concurrently



► Comparison: multi-core vs SMT

- Advantages/disadvantages?

► Comparison: multi-core vs SMT


- Multi-core:
 - Since there are several cores, each is smaller and not as powerful (but also easier to design and manufacture)
 - However, great with thread-level parallelism
- SMT
 - Can have one large and fast superscalar core
 - Great performance on a single thread
 - Mostly still only exploits instruction-level parallelism

The memory hierarchy


- If simultaneous multithreading only:
 - all caches shared
- Multi-core chips:
 - L1 caches private
 - L2 caches private in some architectures and shared in others
- Memory is always shared

Example

- Dual-core Intel Xeon processors
- Each core is hyper-threaded
- Private L1 caches
- Shared L2 caches




Designs with private L2 caches



Both L1 and L2 are private

Examples: AMD Opteron,
AMD Athlon, Intel Pentium D



A design with L3 caches

Example: Intel Itanium 2 32

Private vs shared caches?

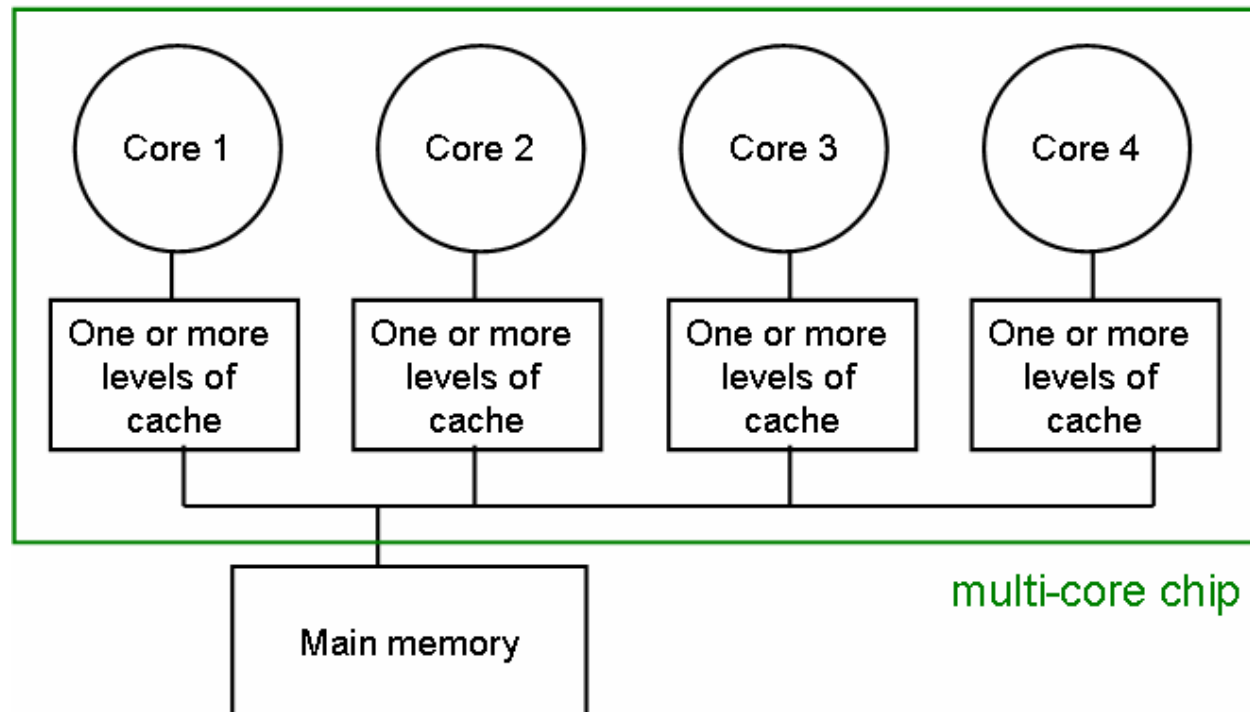
- Advantages/disadvantages?

Private vs shared caches

- Advantages of private:
 - They are closer to core, so faster access
 - Reduces contention
- Advantages of shared:
 - Threads on different cores can share the same cache data
 - More cache space available if a single (or a few) high-performance thread runs on the system

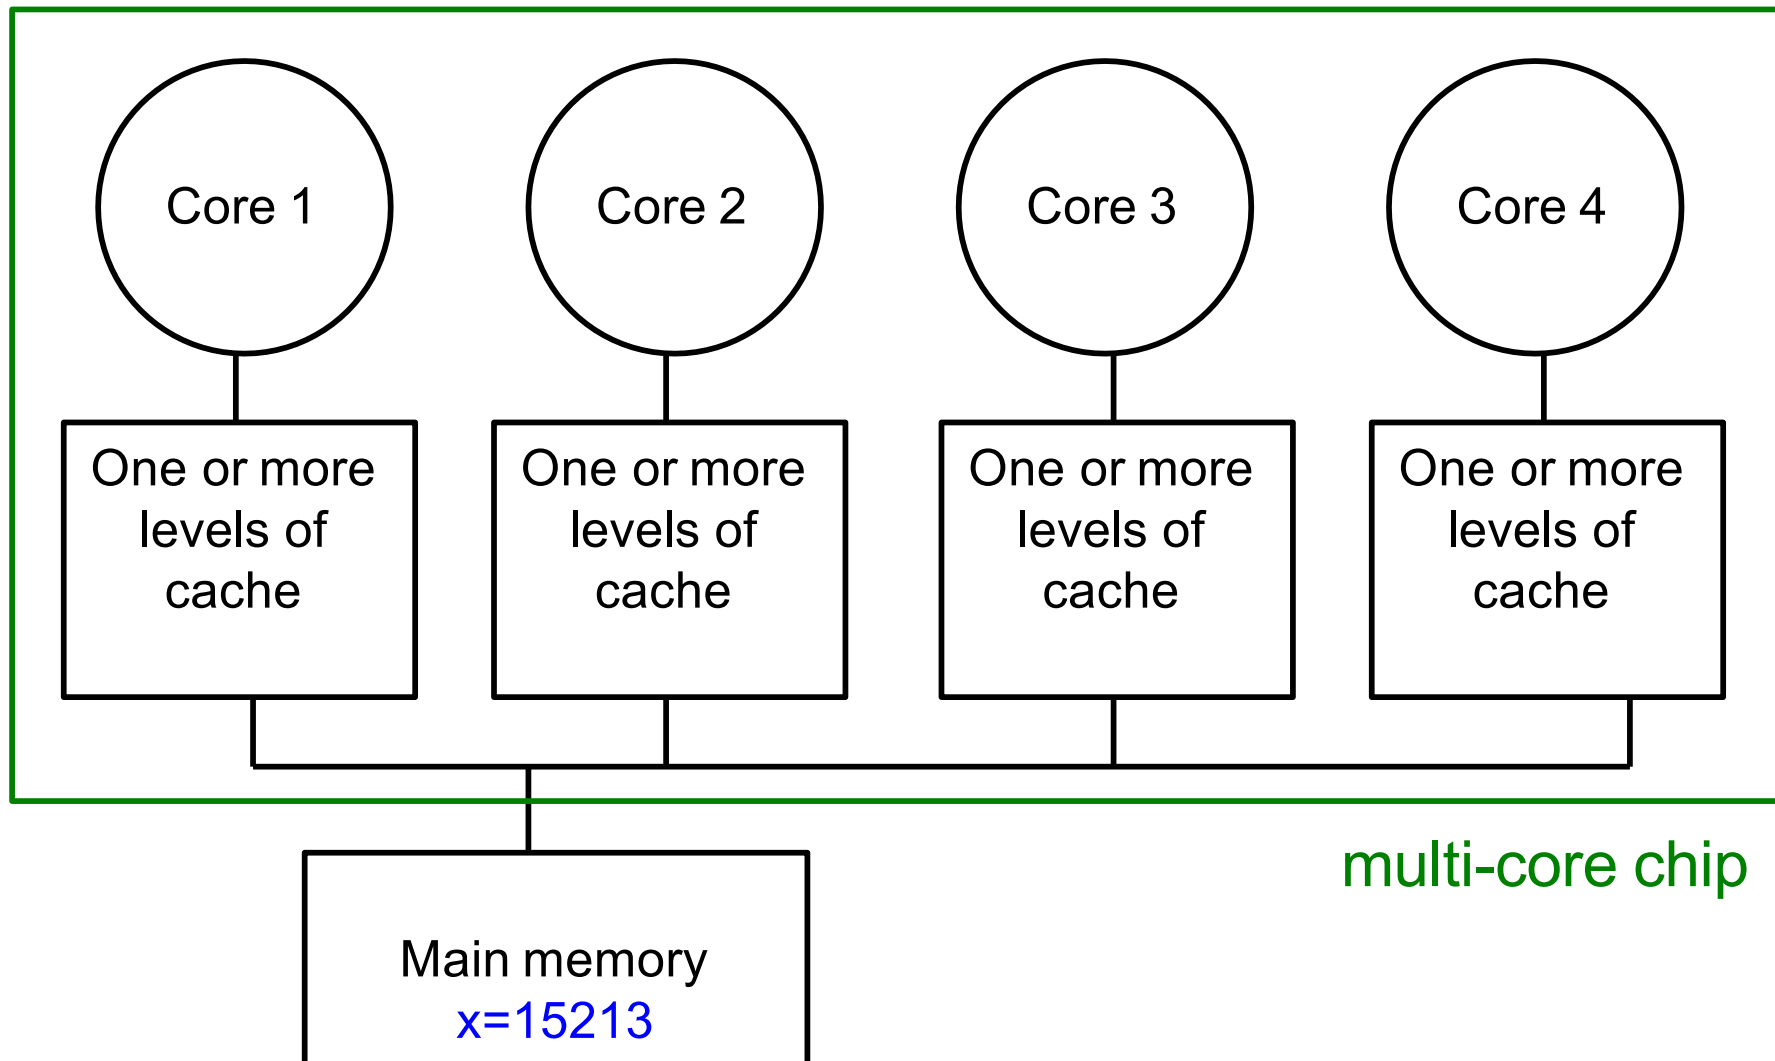
► The cache coherence problem

- Since we have private caches:
How to keep the data consistent across caches?
- Each core should perceive the memory as a monolithic array, shared by all the cores



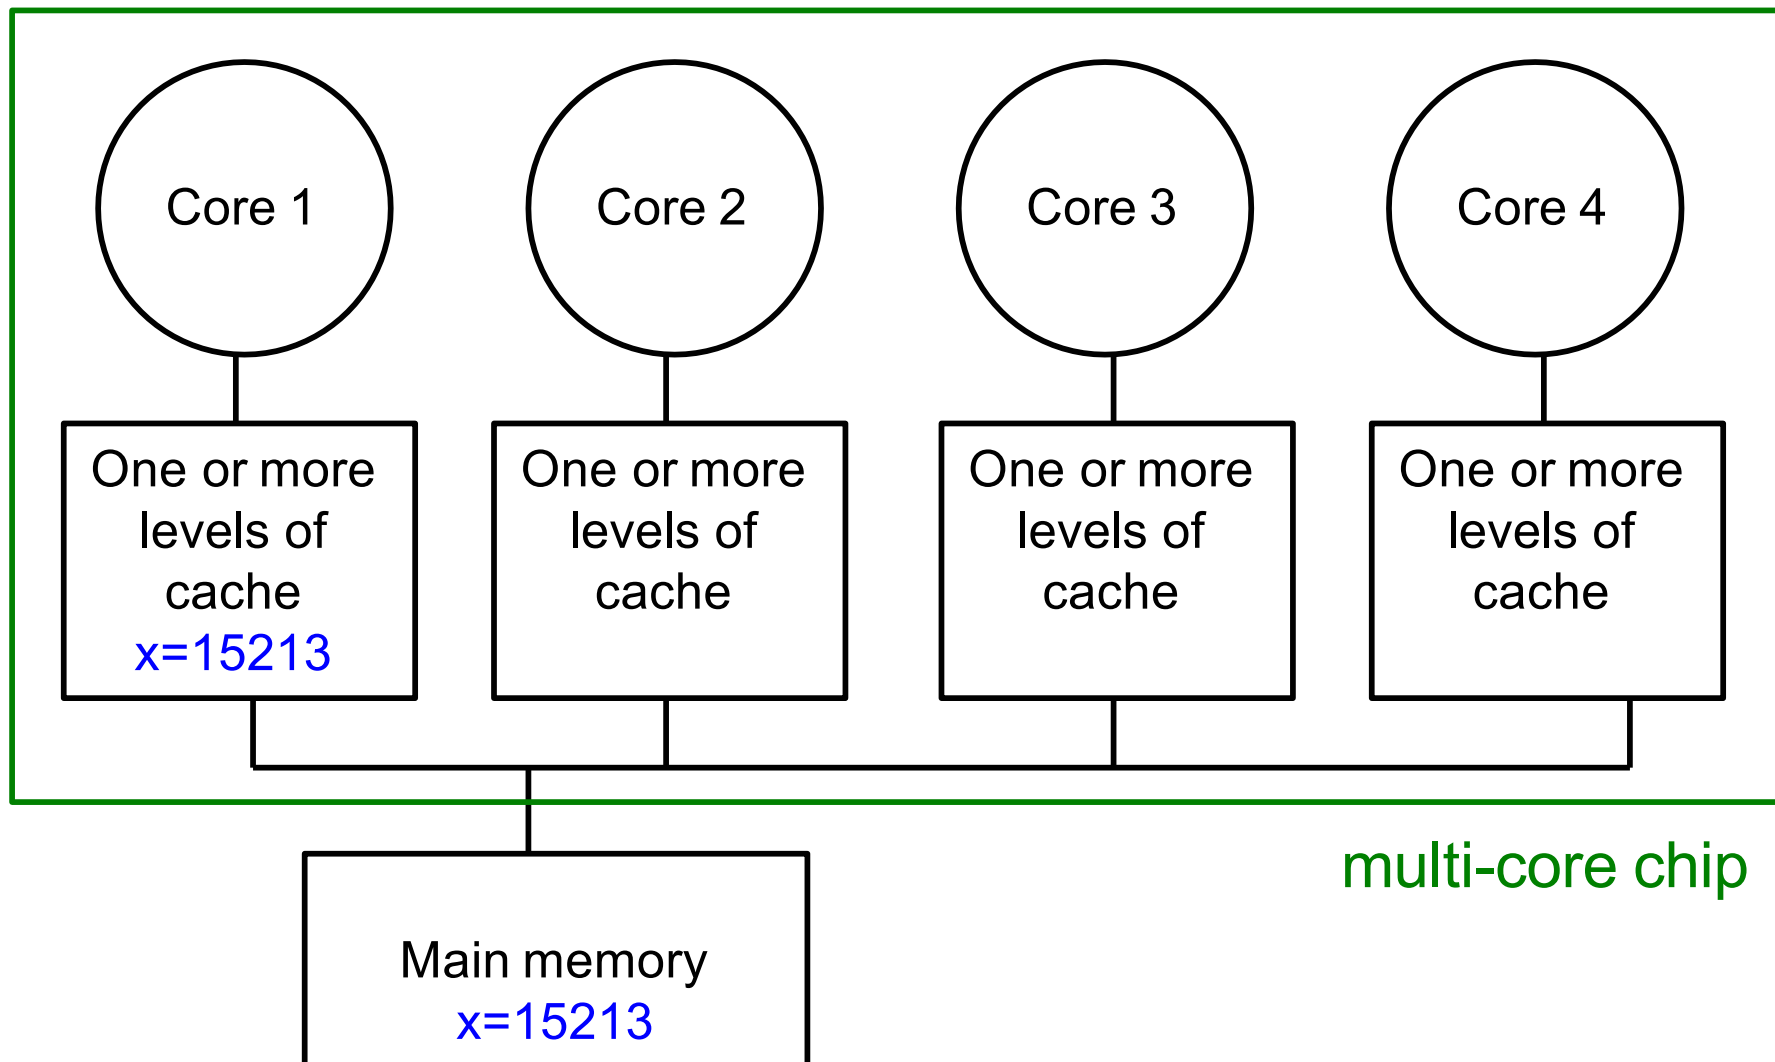
The cache coherence problem

Suppose variable x initially contains 15213



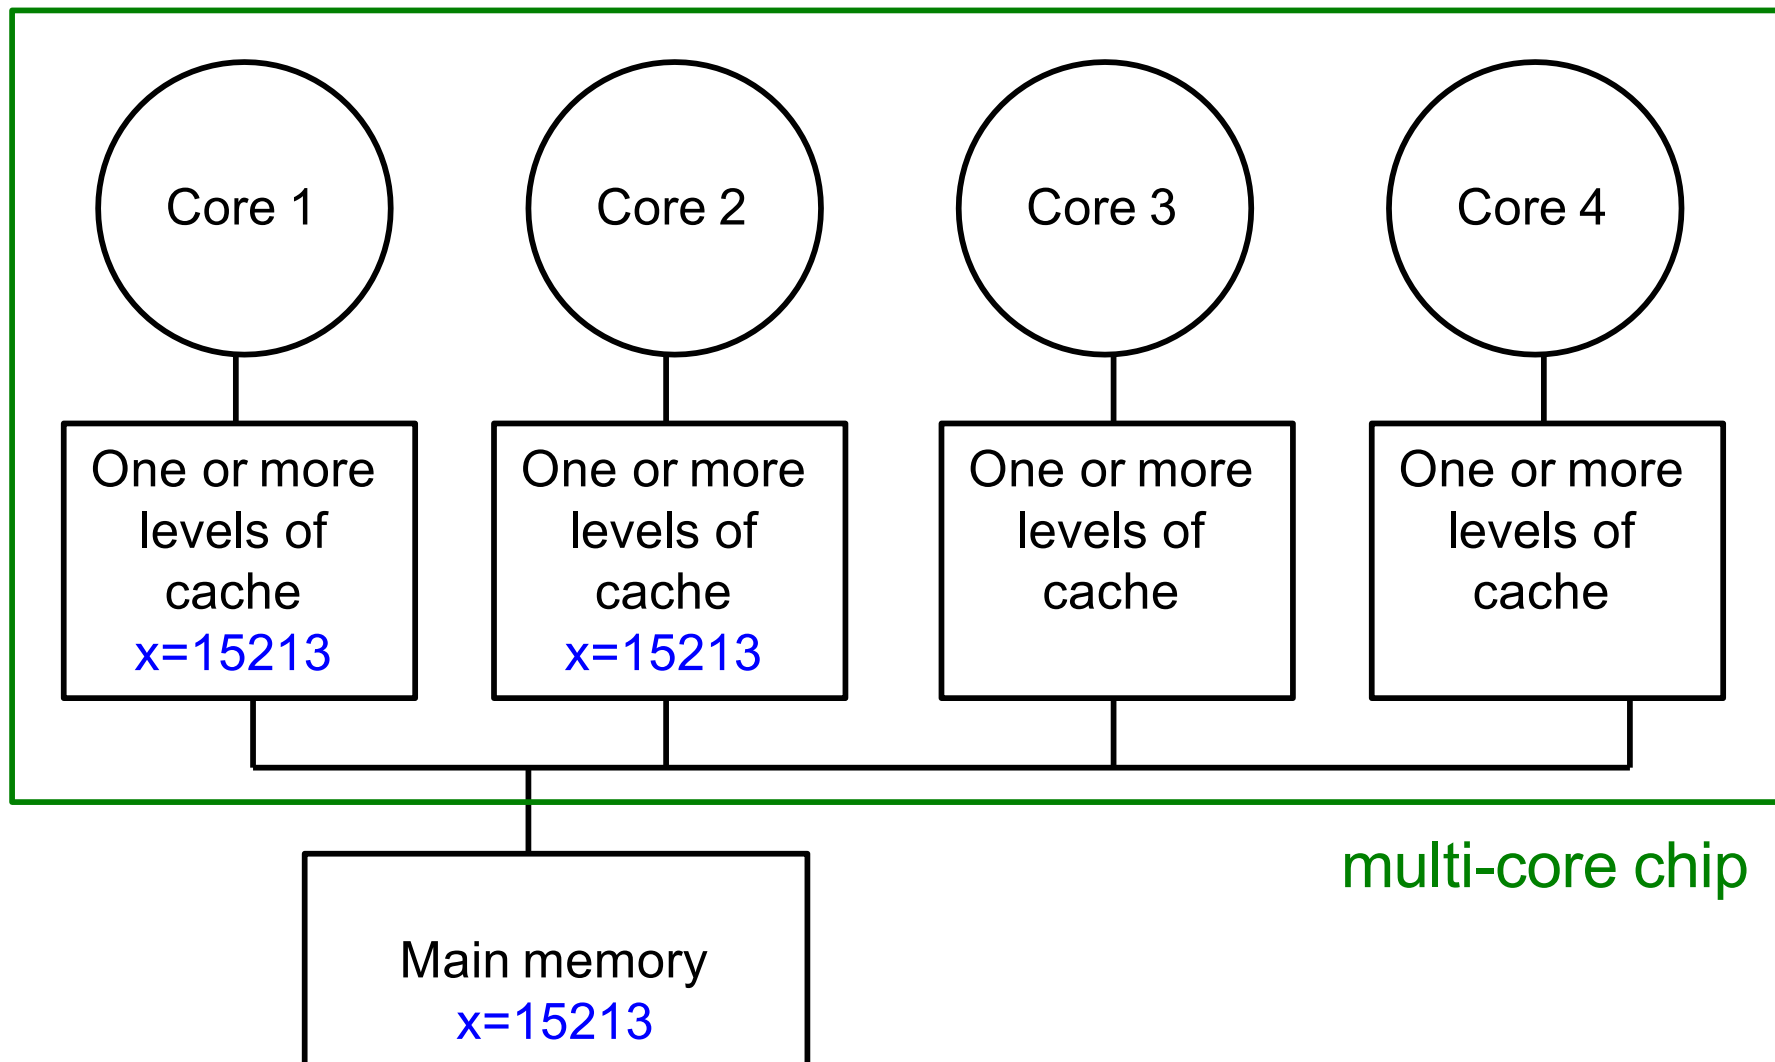
The cache coherence problem

Core 1 reads x



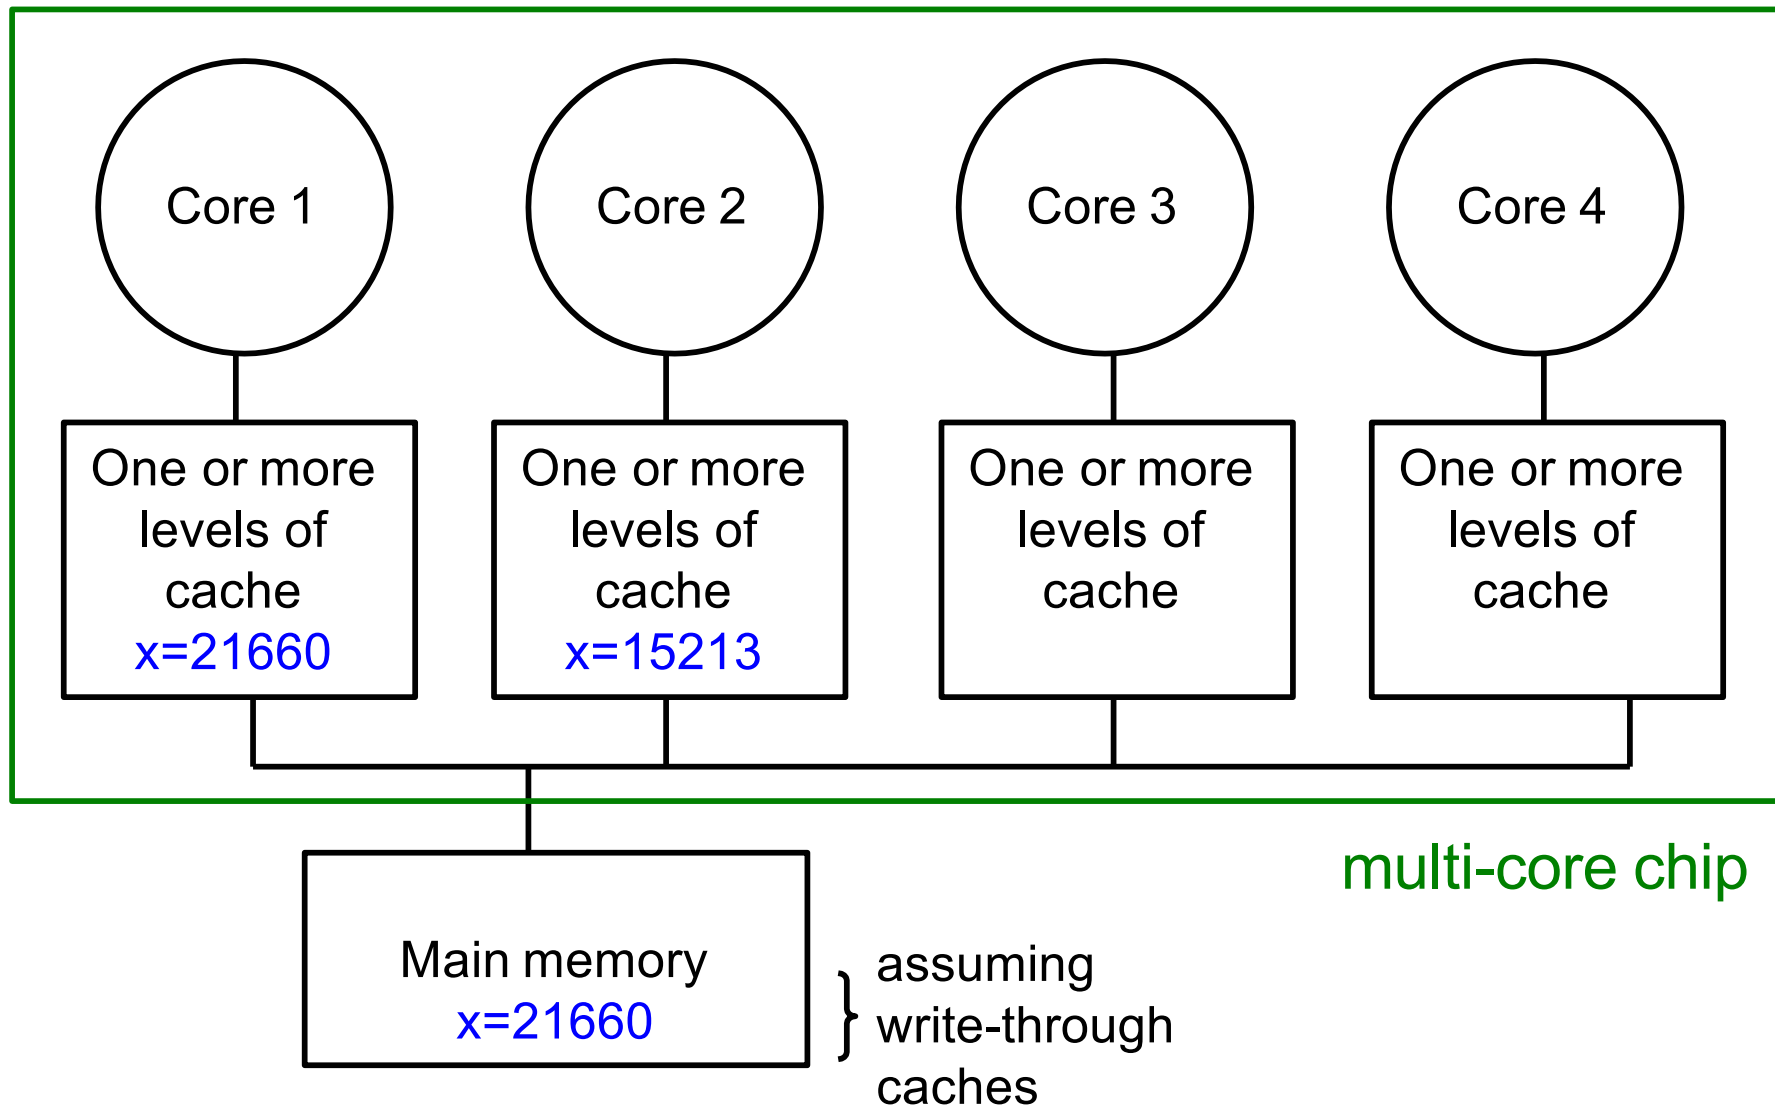
The cache coherence problem

Core 2 reads x



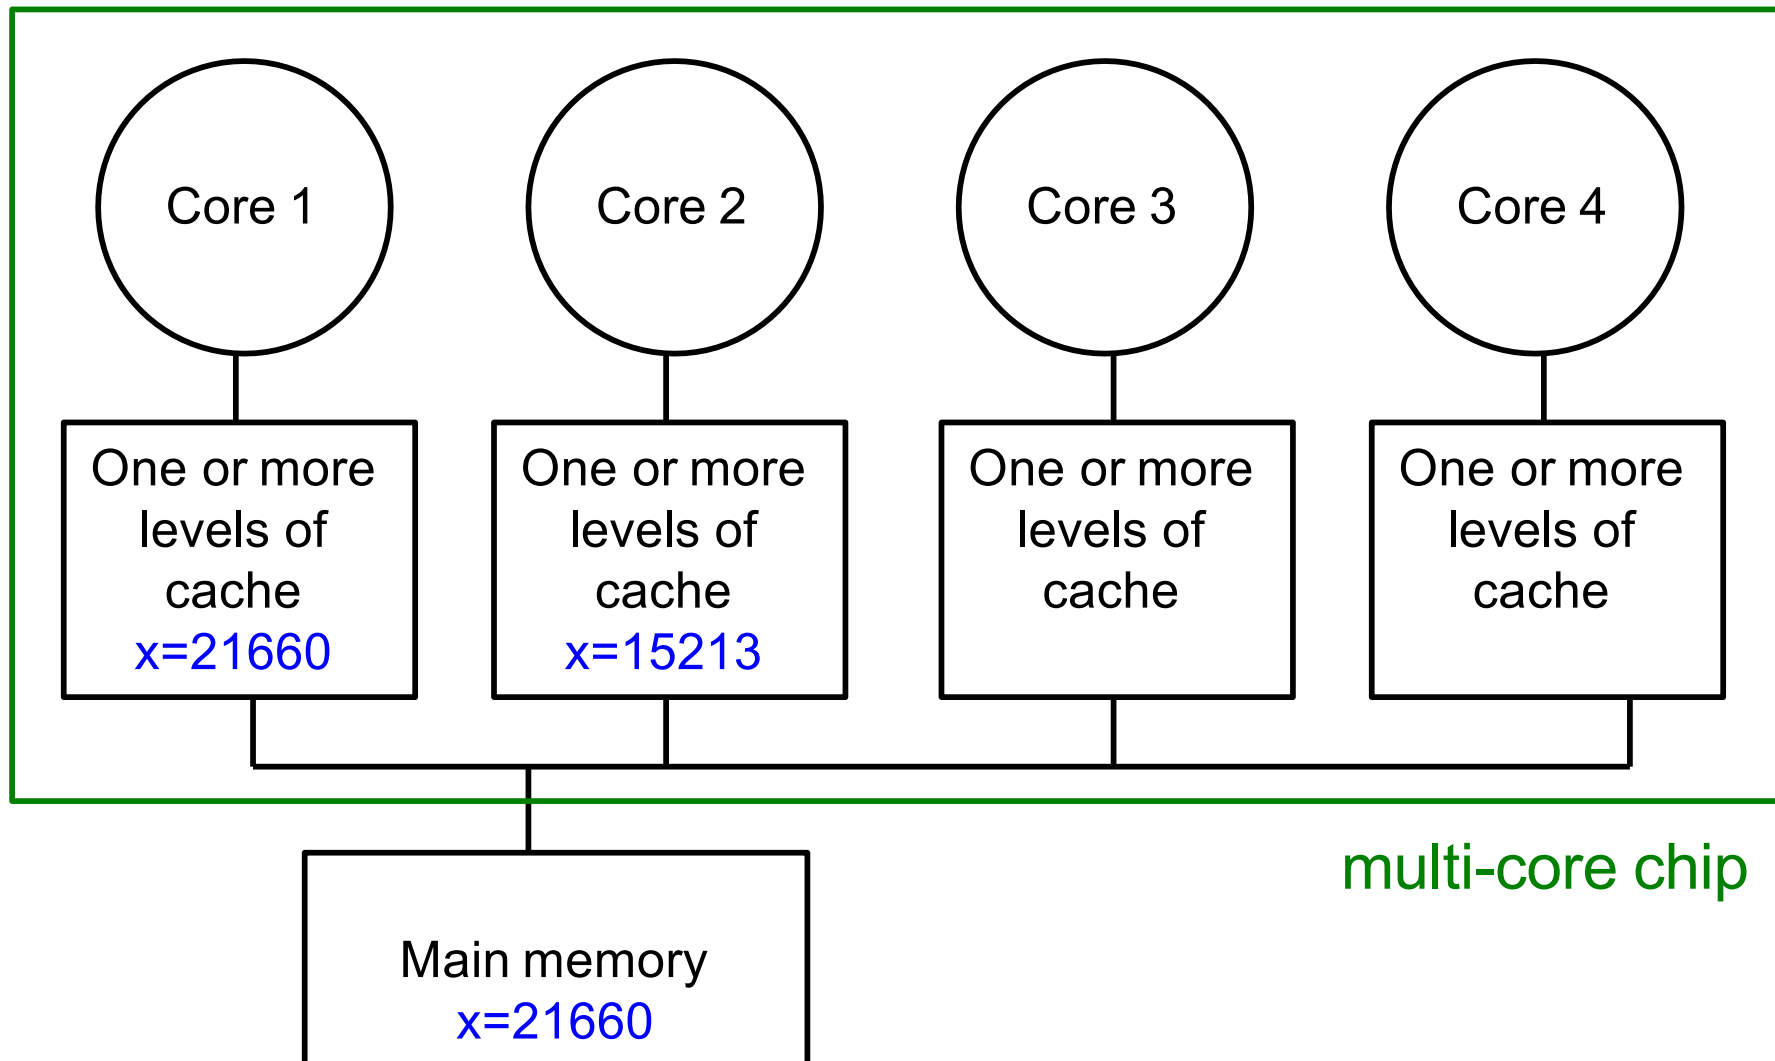
The cache coherence problem

Core 1 writes to x , setting it to 21660



The cache coherence problem

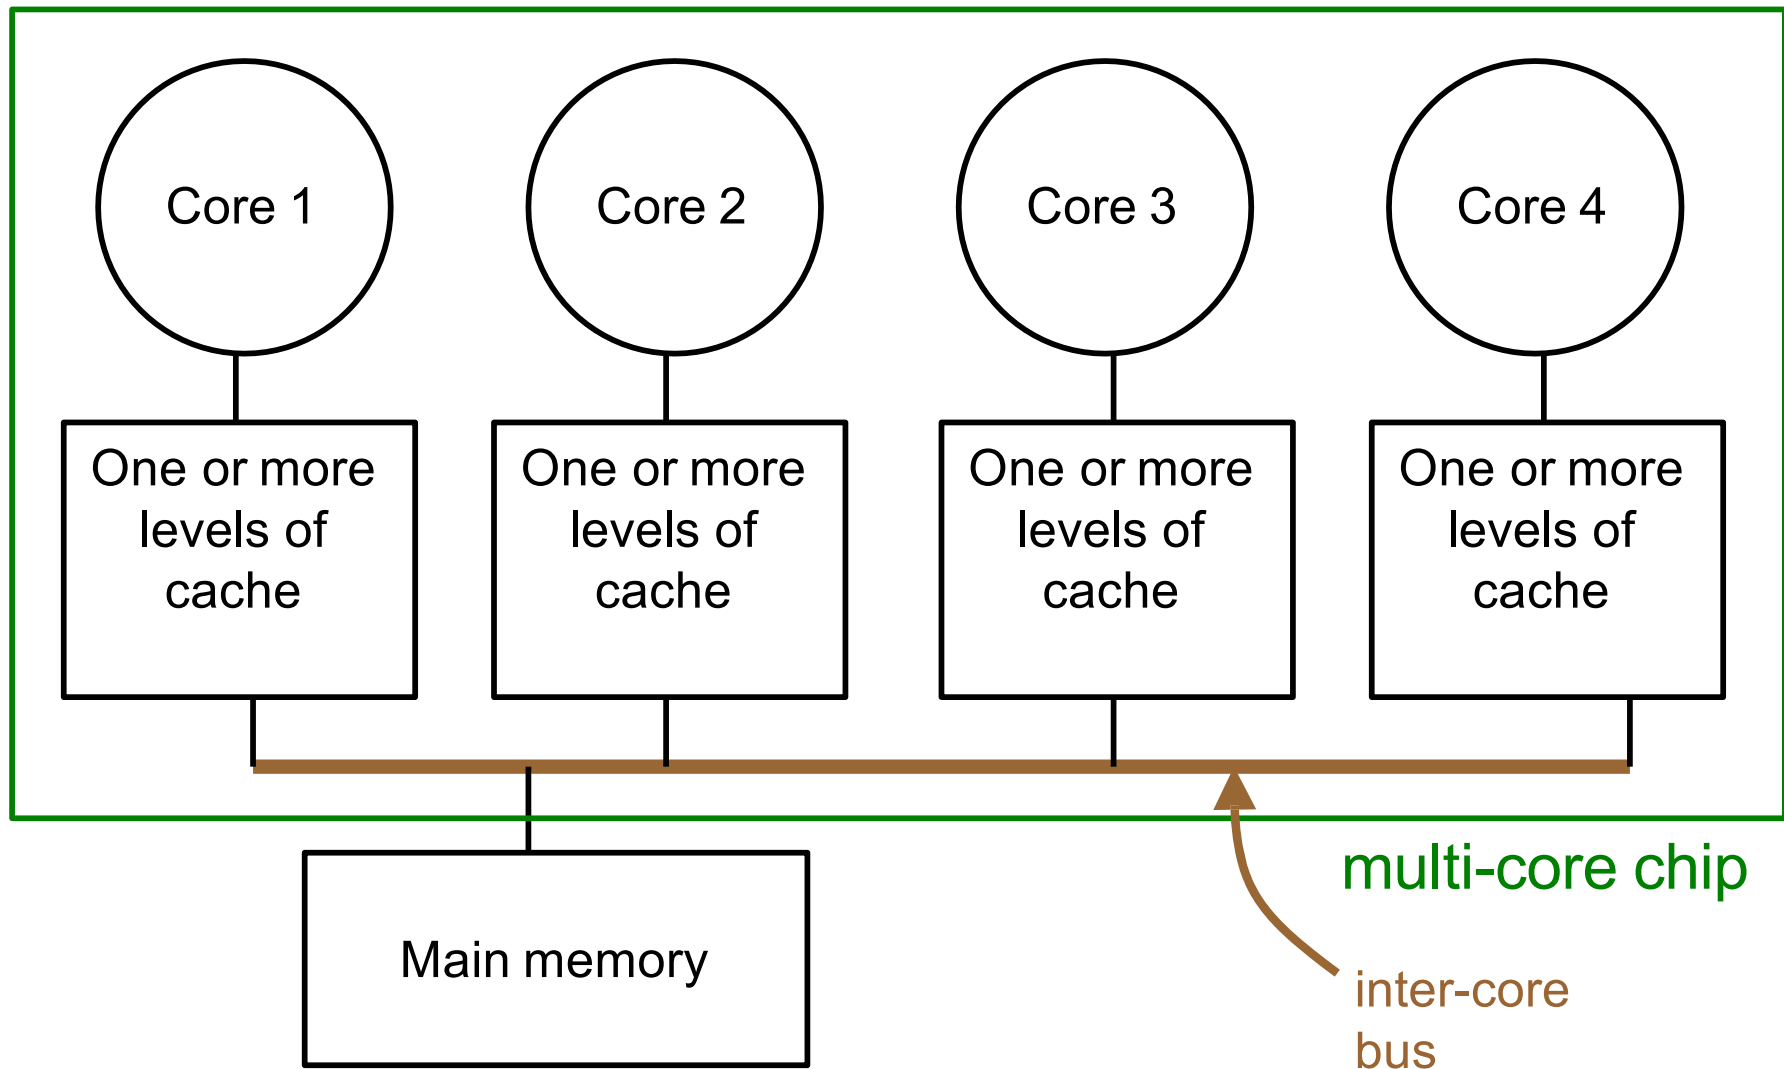
Core 2 attempts to read x ... gets a stale copy



Solutions for cache coherence

- This is a general problem with multiprocessors, not limited just to multi-core
- There exist many solution algorithms, coherence protocols, etc.
- A simple solution:
invalidation-based protocol with snooping

Inter-core bus

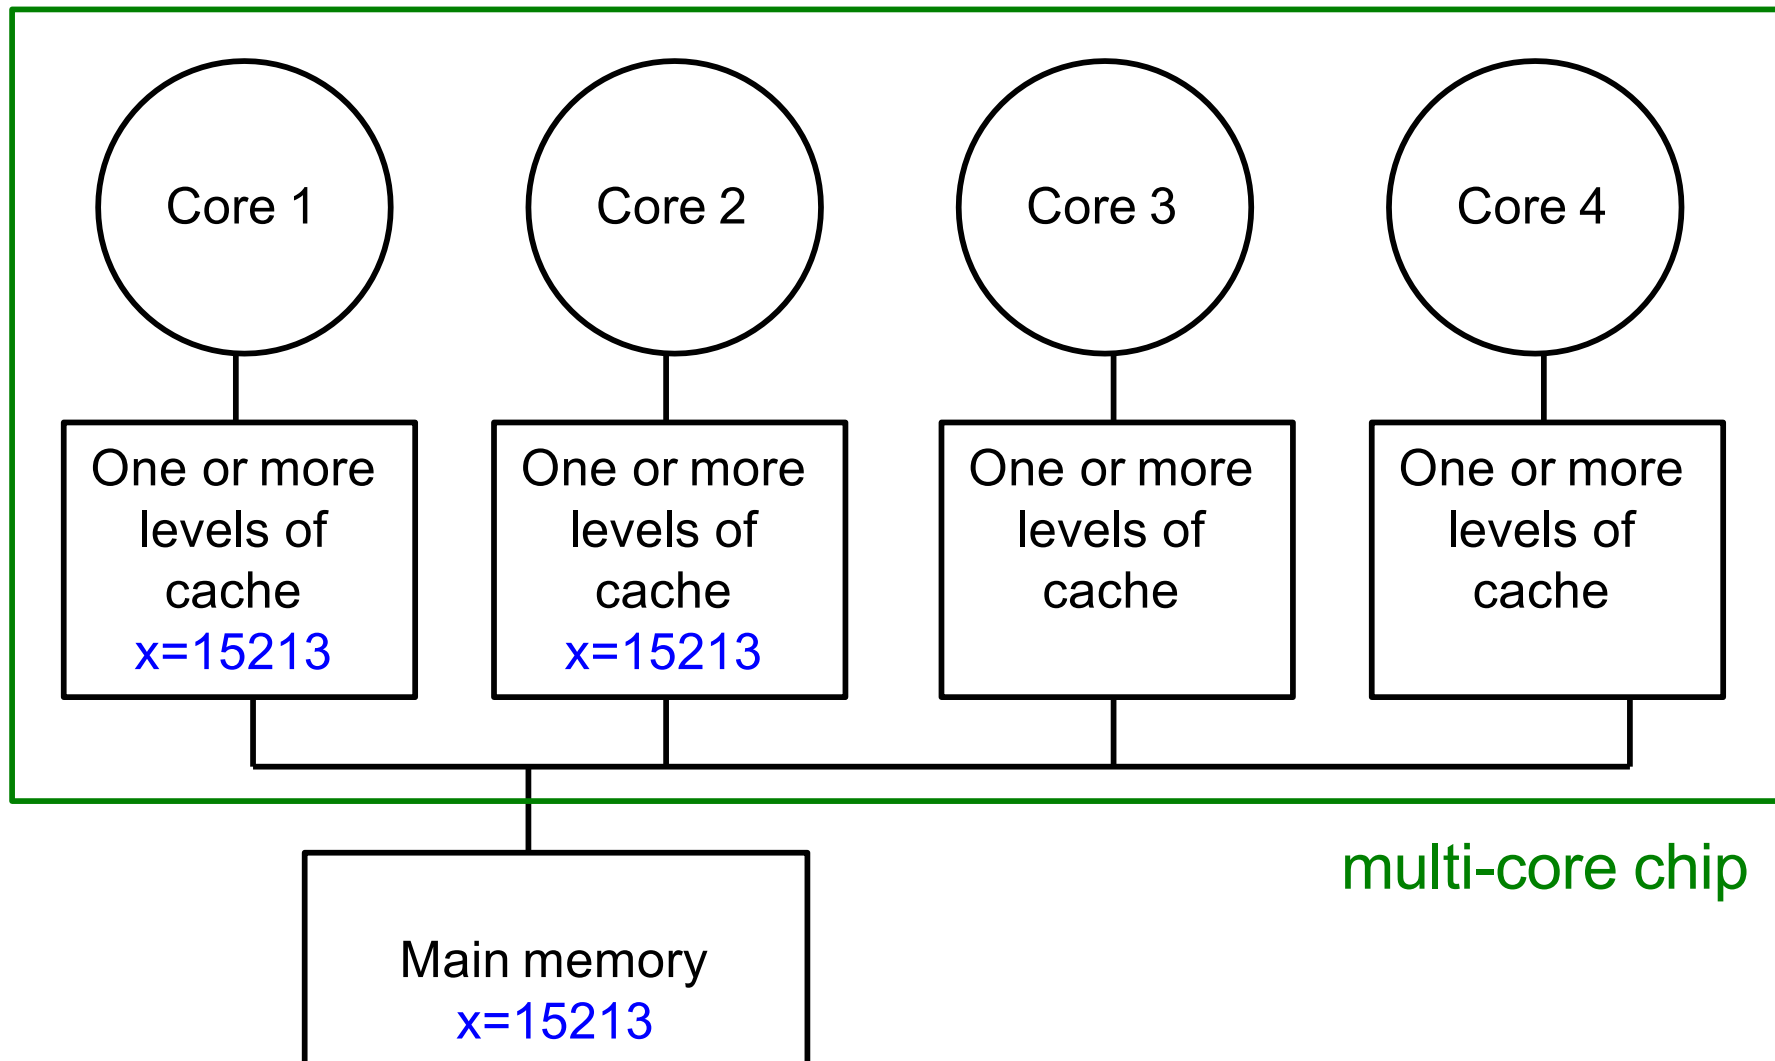


Invalidation protocol with snooping

- Invalidation:
If a core writes to a data item, all other copies of this data item in other caches are *invalidated*
- Snooping:
All cores continuously “snoop” (monitor) the bus connecting the cores.

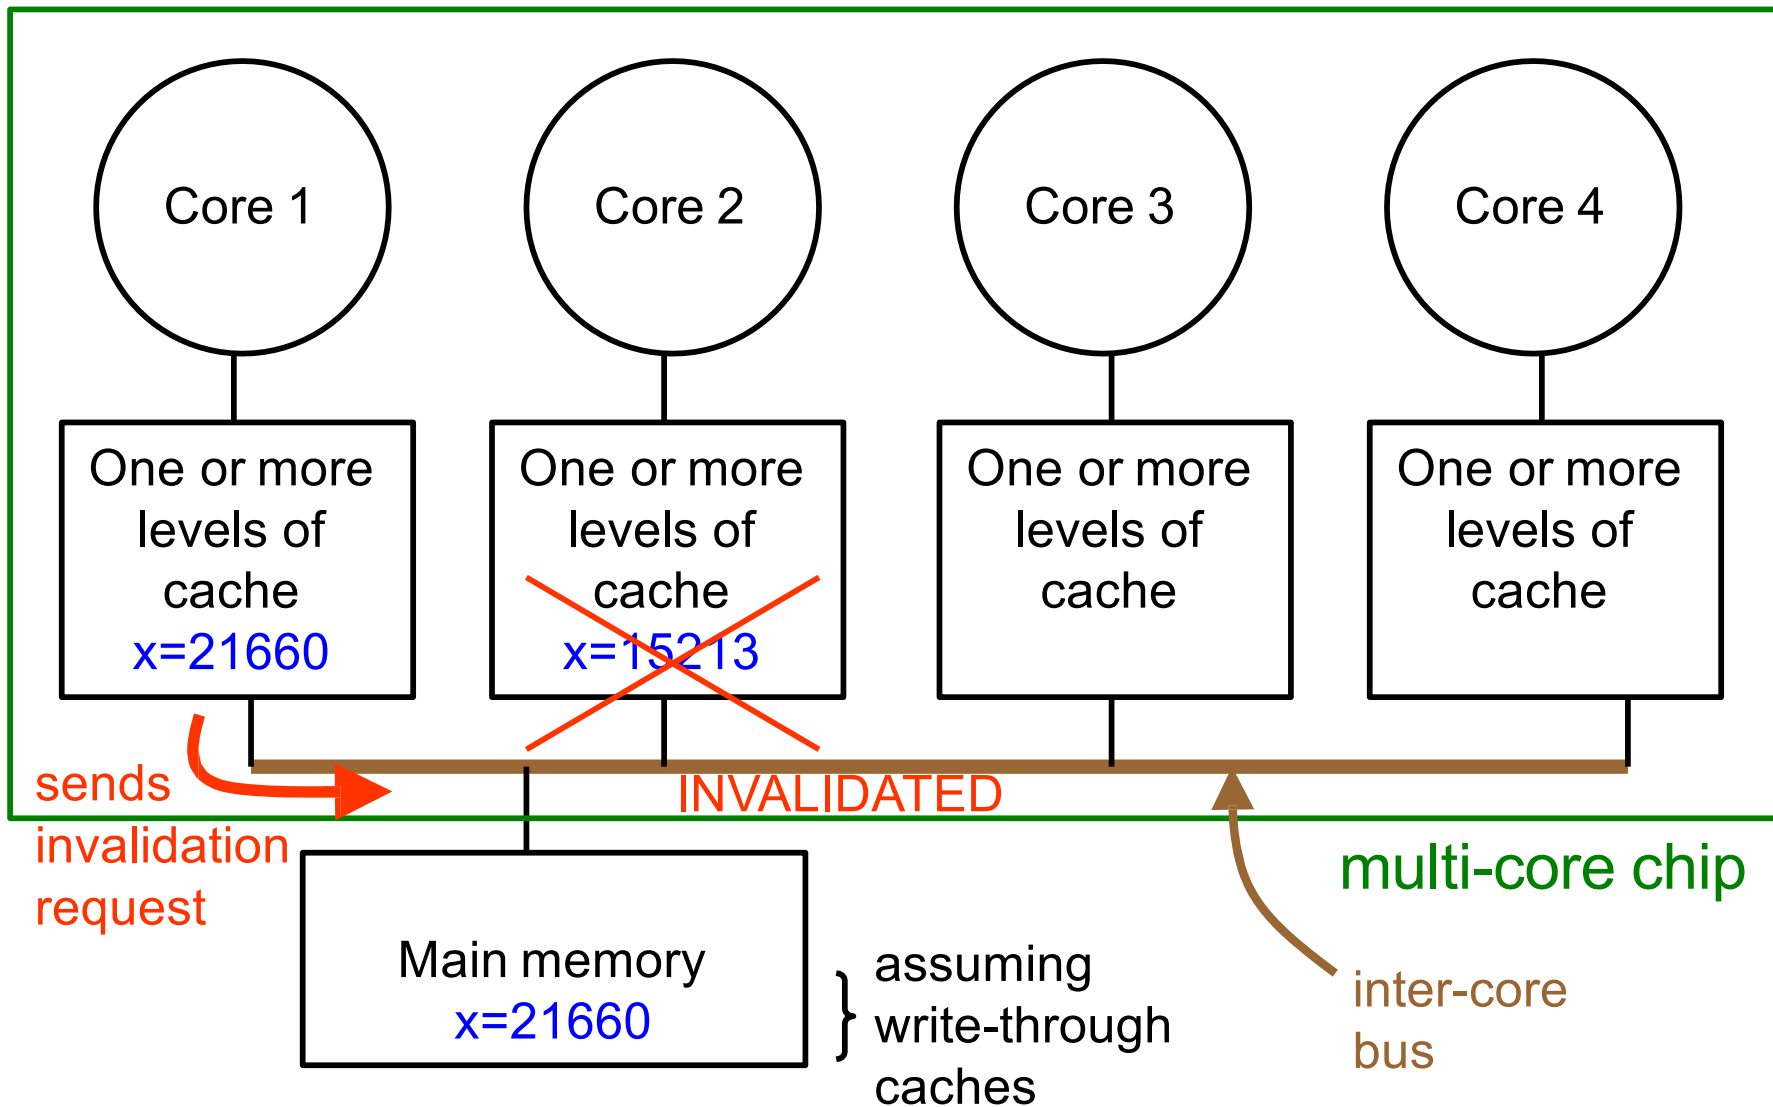
The cache coherence problem

Revisited: Cores 1 and 2 have both read x



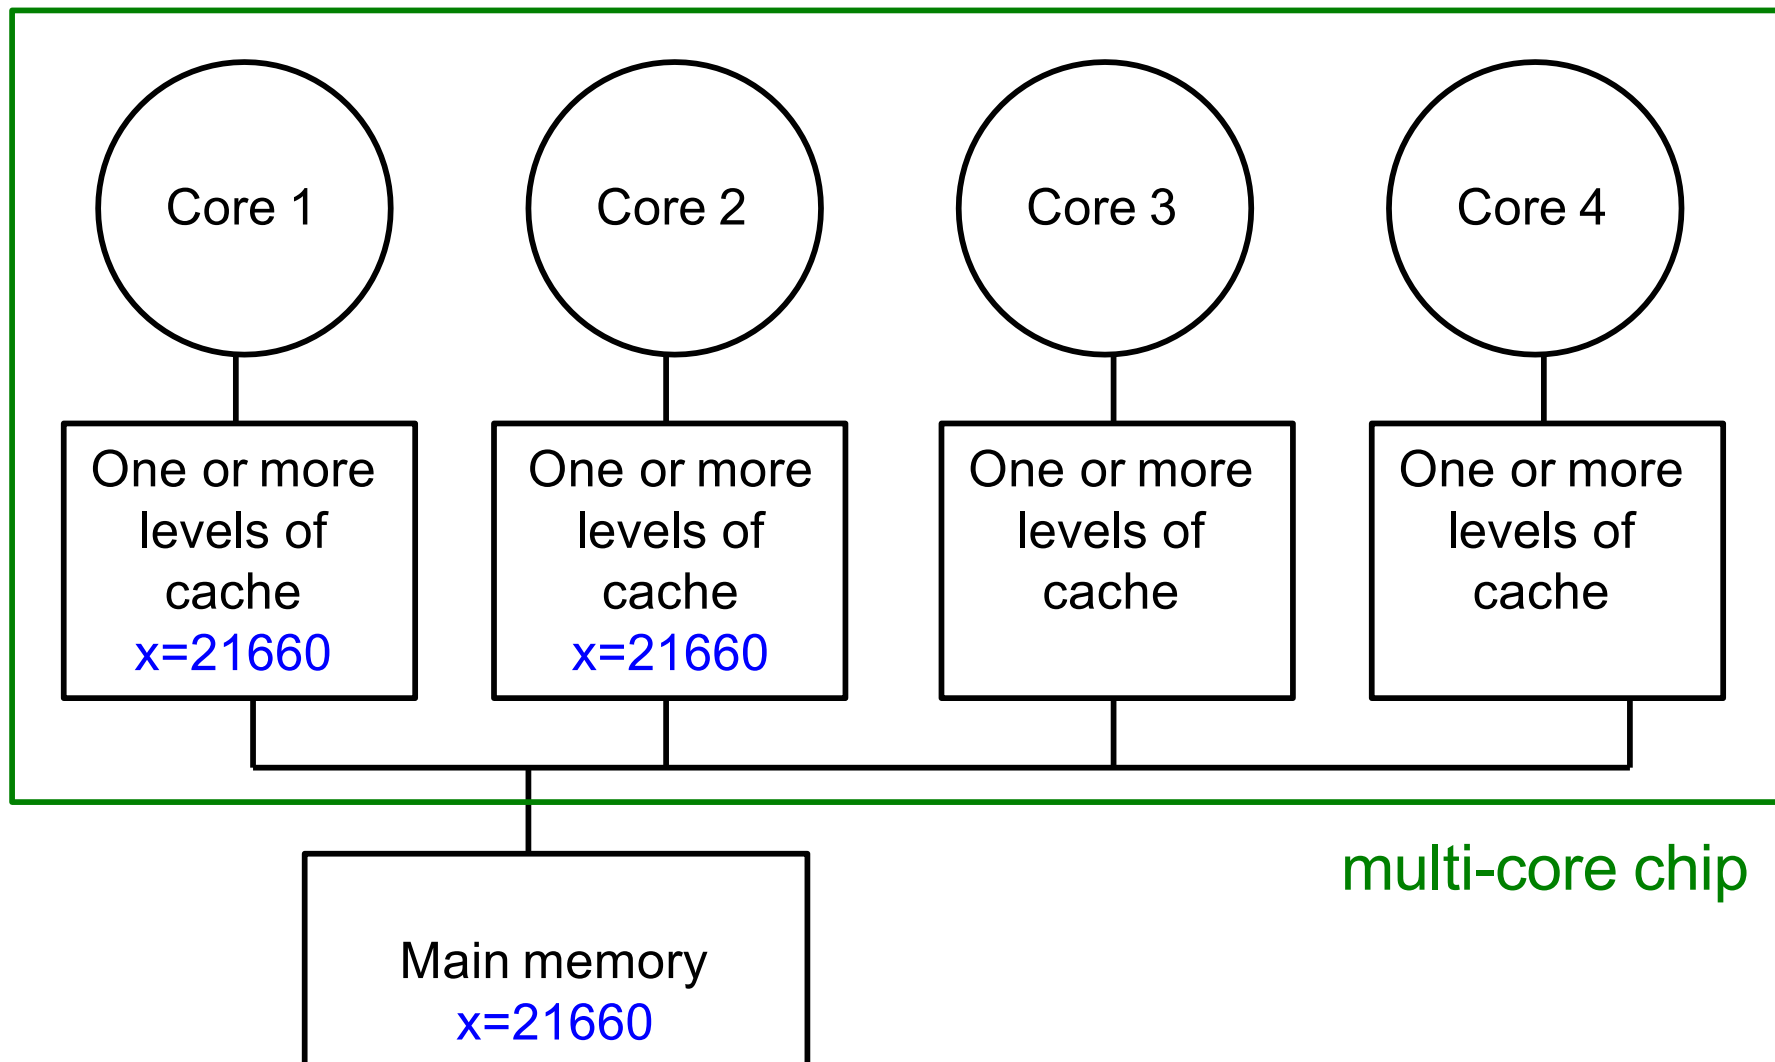
The cache coherence problem

Core 1 writes to x, setting it to 21660



The cache coherence problem

Core 2 reads x . Cache misses, and loads the new copy.



▶ Programming for multi-core

- Programmers must use threads or processes
- Spread the workload across multiple cores
- Write parallel algorithms
- OS will map threads/processes to cores

Thread safety very important

- Pre-emptive context switching:
context switch can happen AT ANY TIME
- True concurrency, not just uniprocessor time-slicing
- Concurrency bugs exposed much faster with multi-core

► **However: Need to use synchronization**

even if only time-slicing on a uniprocessor

```
int counter=0;
```

```
void thread1() {  
    int temp1=counter;  
    counter = temp1 + 1;  
}
```

```
void thread2() {  
    int temp2=counter;  
    counter = temp2 + 1;  
}
```

Need to use synchronization even if only time-slicing on a uniprocessor

```
temp1=counter;  
counter = temp1 + 1;  
temp2=counter;  
counter = temp2 + 1
```

gives counter=2

```
temp1=counter;  
temp2=counter;  
counter = temp1 + 1;  
counter = temp2 + 1
```


gives counter=1
or 2

Assigning threads to the cores

- Each thread has an *affinity mask*
- Affinity mask specifies what cores the thread is allowed to run on
- Different threads can have different masks
- Affinities are inherited across `fork()`

► Affinity masks are bit vectors


- Example: 4-way multi-core, without SMT



- Process/thread is allowed to run on cores 0,2,3, but not on core 1

Affinity masks when multi-core and SMT combined

- Separate bits for each simultaneous thread
- Example: 4-way multi-core, 2 threads per core



- Core 2 can't run the process
- Core 1 can only use one simultaneous thread

Default Affinities

- Default affinity mask is all 1s:
all threads can run on all processors
- Then, the OS scheduler decides what threads run on what core
- OS scheduler detects skewed workloads, migrating threads to less busy processors

Process migration is costly

- Need to restart the execution pipeline
- Cached data is invalidated
- OS scheduler tries to avoid migration as much as possible:
it tends to keep a thread on the same core
- This is called *soft affinity*

Hard affinities

- The programmer can prescribe her own affinities (hard affinities)
- Rule of thumb: use the default scheduler unless a good reason not to

▶ When to set your own affinities

- Two (or more) threads share data-structures in memory
 - map to same core so that can share cache
- Real-time threads:
Example: a thread running a robot controller:
 - must not be context switched, or else robot can go unstable
 - dedicate an entire core just to this thread



Source: Sensable.com

Kernel scheduler API

```
#include <sched.h>

int sched_getaffinity(pid_t pid,
    unsigned int len, unsigned long * mask);
```

Retrieves the current affinity mask of process 'pid' and stores it into space pointed to by 'mask'.

'len' is the system word size: sizeof(unsigned int long)

Kernel scheduler API

```
#include <sched.h>

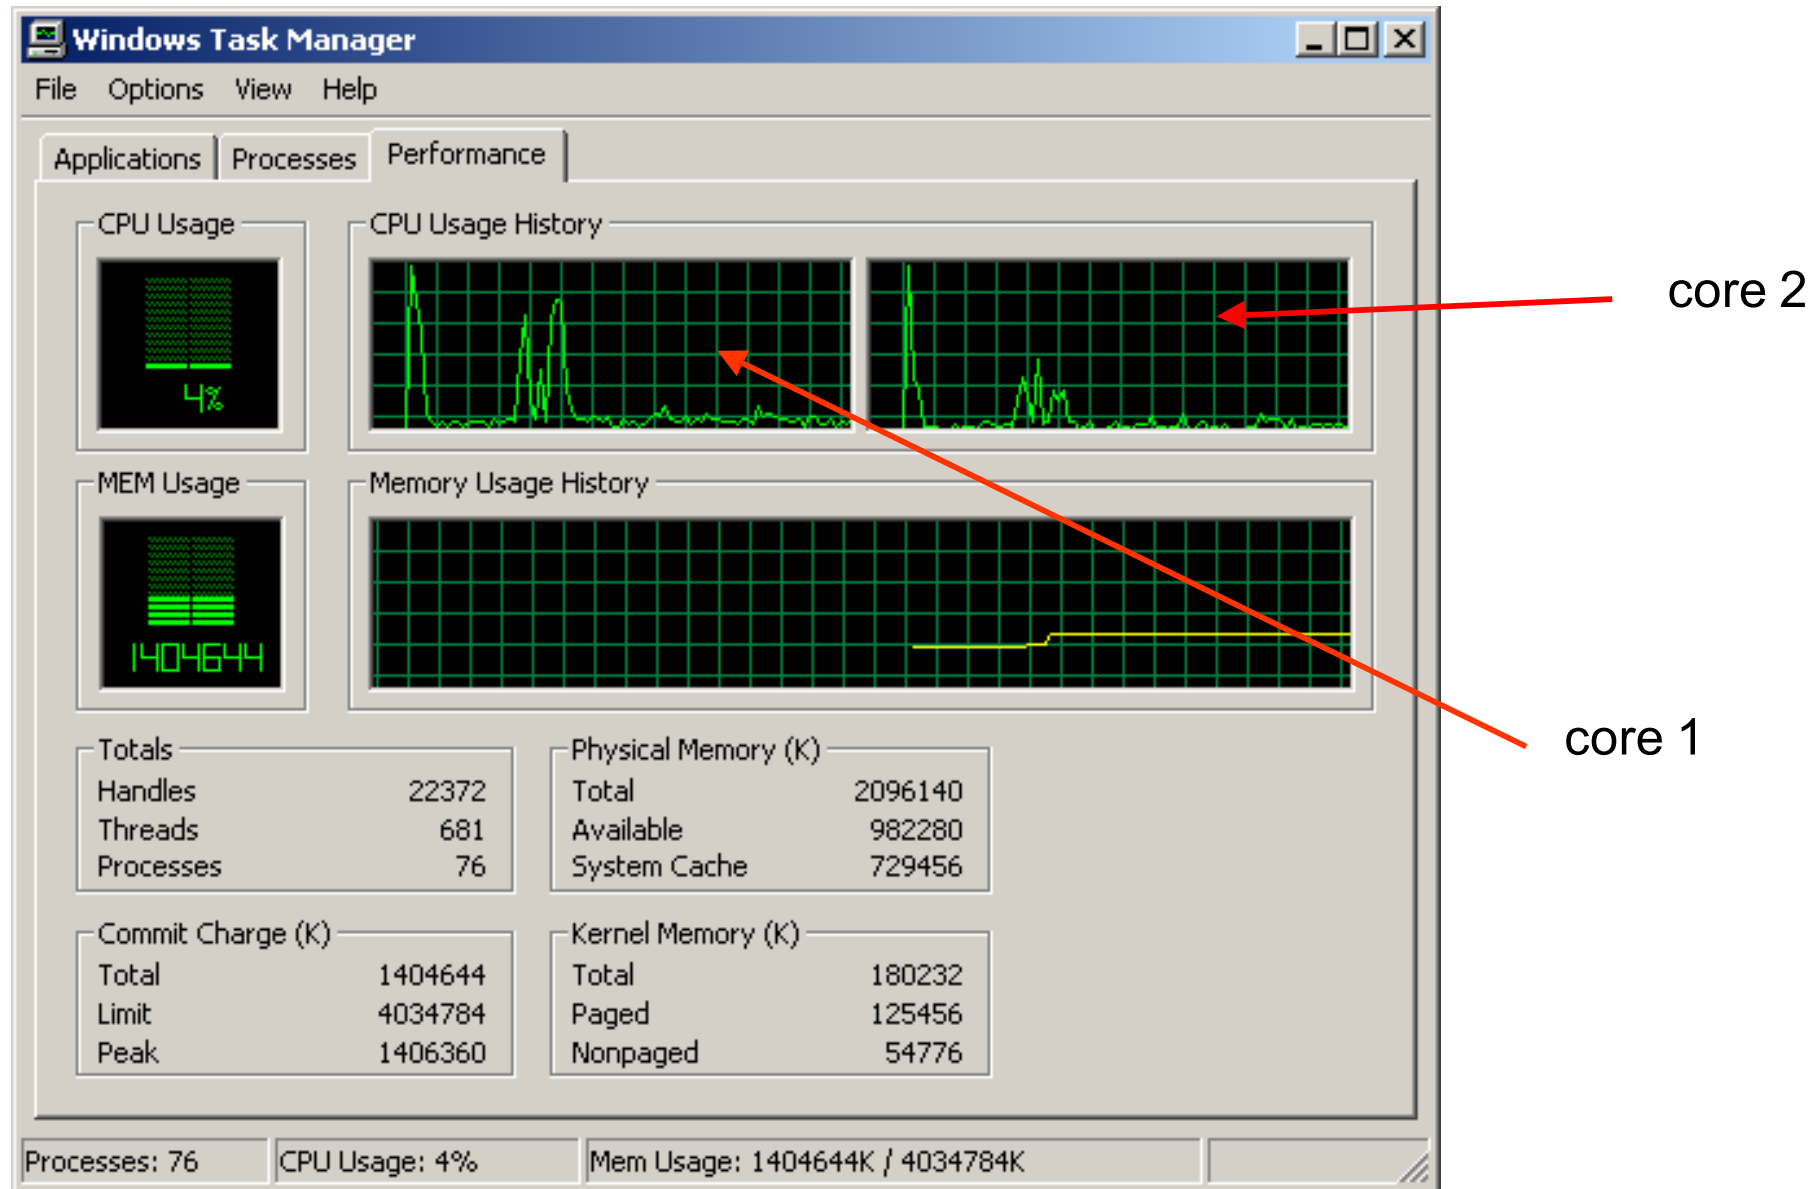
int sched_setaffinity(pid_t pid,
    unsigned int len, unsigned long * mask);
```

Sets the current affinity mask of process 'pid' to *mask
'len' is the system word size: sizeof(unsigned int long)

To query affinity of a running process:


```
[~]$ taskset -p 3935 pid 3935's
current affinity mask: f
```

Windows Task Manager



Conclusion

- Multi-core chips an important new trend in computer architecture
- Several new multi-core chips in design phases
- Parallel programming techniques likely to gain importance



► Further readings

A. S. Tanenbaum, T. Austin, and B. R. Chandavarkar, *Structured computer organization*, 6. ed ., international ed. Boston, Mass.: Pearson, 2013 – Chapter 8

On-Chip Paralellism

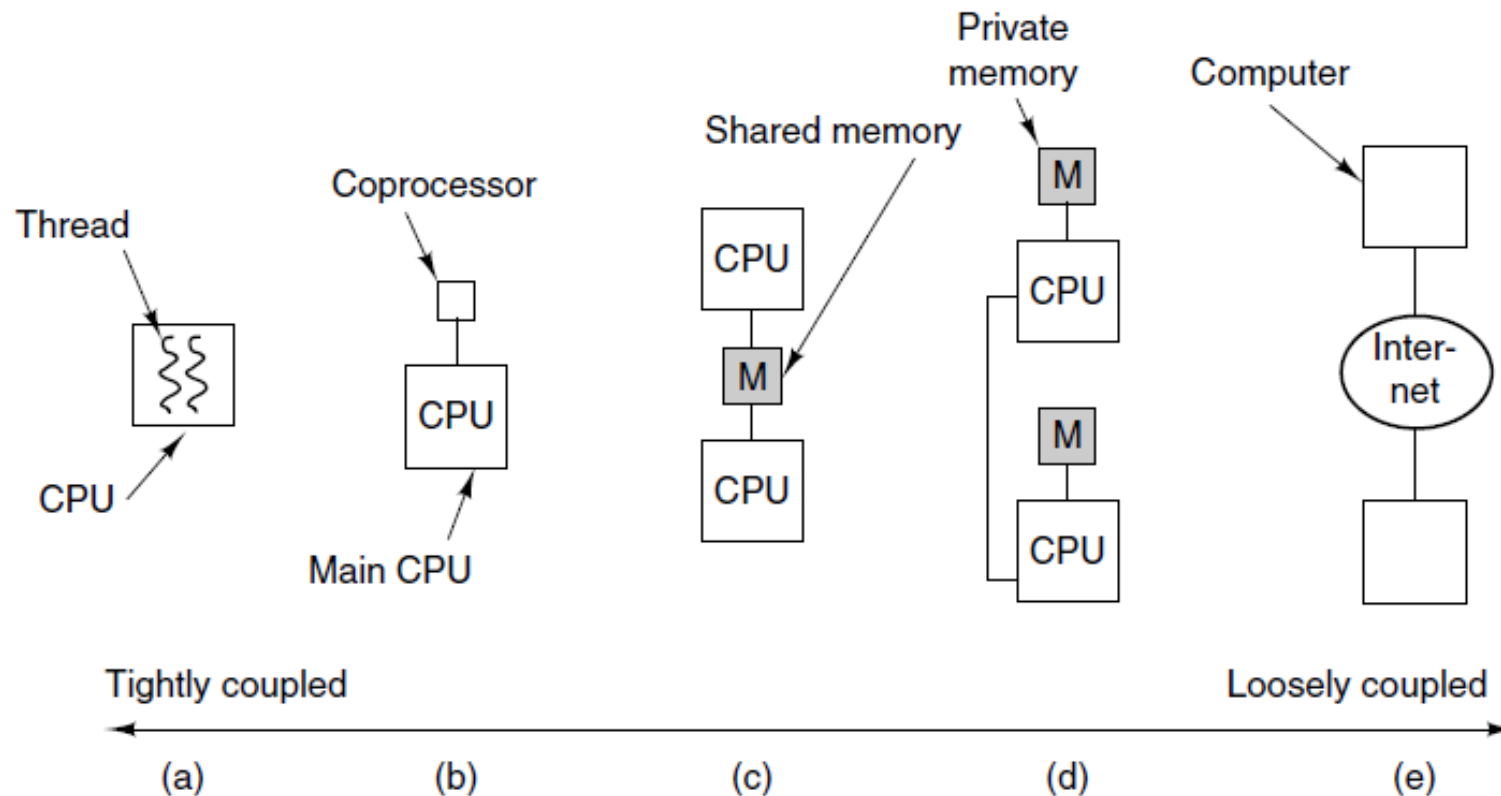


Figure 8-1. (a) On-chip parallelism. (b) A coprocessor. (c) A multiprocessor. (d) A multicomputer. (e) A grid.


► Programming

C/C++ Threads – short recap

Sequential program

- A sequential program can be presented as a control flow graph


```
0:      public static int manhattan(int a, int b){  
1:          if (a < 0)  
2:              a = -a;  
3:          if (b < 0)  
4:              b = -b;  
5:          return a+b;  
      }
```



Thread


- For a specific input
 - a path of the control graph is run through

`manhattan(-1, 1)`



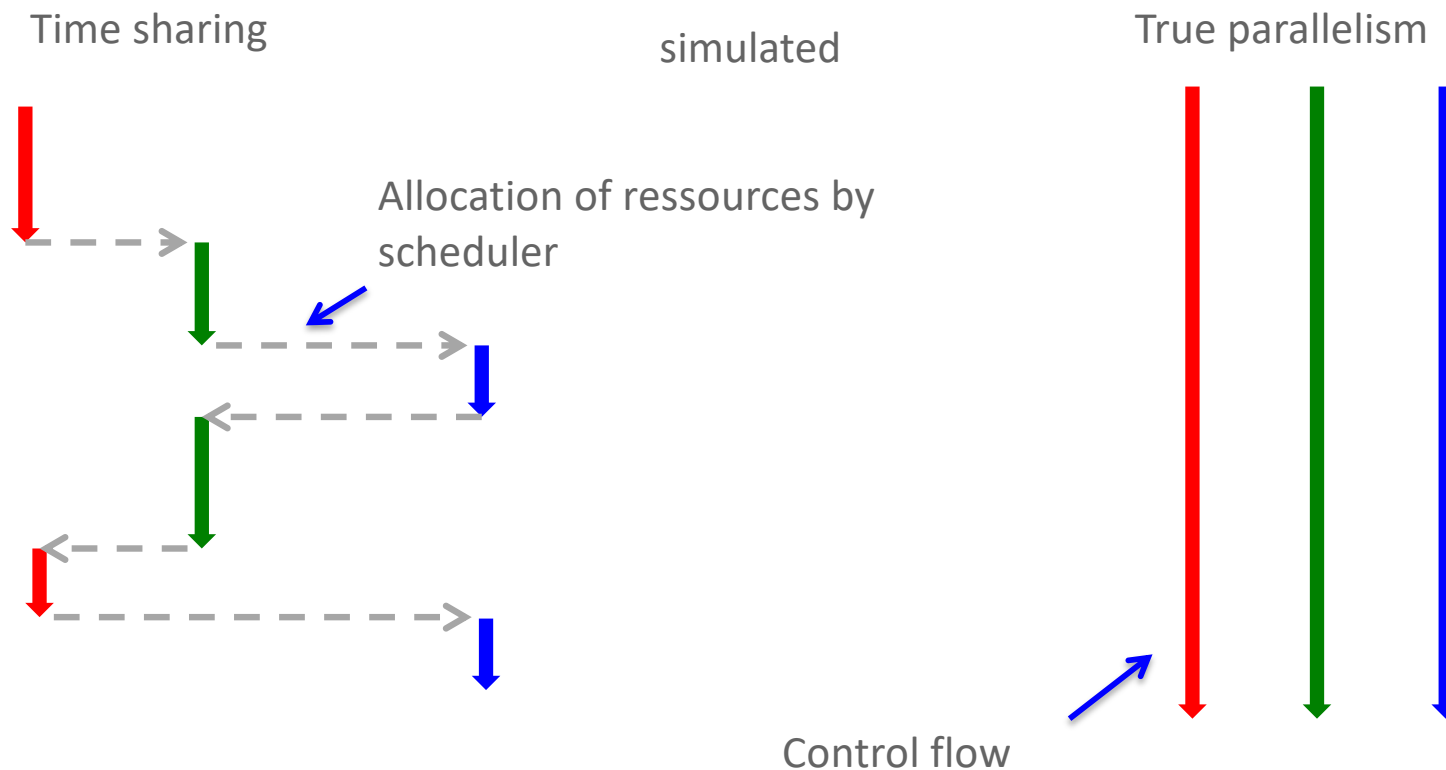
Concurrent

- Program (process) are called concurrent if they do not depend on each other
 - The associated control flows can be run through in parallel



Concurrent

- Concurrent control flows can be processed in parallel if the required resources (CPU, printer, ...) are available for each control flow.
- Otherwise, the resources are temporarily allocated to the control flow by a scheduler.



► 71 **Concurrent**

Typical applications for concurrent programming


- Concurrency is given by the task in a natural way (e.g. motor control)
- Better utilization of resources (hardware)
 - Modern processors typically have multiple CPU cores
- Increasing the throughput of a system
 - Time-consuming I/O request blocks only one control flow and not the entire application
- Avoidance of blockades
 - Improvement of the interactivity of user interfaces
 - Reduction of the influence of blocking interfaces (e.g. network connections)

- A process is an executed program
 - A process receives resources from the operating system for a certain period of time.
 - Processor
 - Own main memory area (address space)

- A thread is a lightweight process
 - Belongs to a process
 - Shares address space (data) and resources with other threads of the same process
 - Has its own register set and stack (e.g. for local variables)
 - can be stopped and continued by a scheduler
 - can wait for an event without CPU load

Processes vs. threads

Operating system



- Threads typically require less administrative effort
 - The scheduler can switch between threads faster
 - Data exchange between threads is accelerated by the common address space.
 - but also holds dangers...

- Application area of threads: subtasks of a program should be executed in parallel
 - GUI
 - technical logic
 - DB and network accesses

► Multithreaded communication

► Serial execution:

- All our programs so far has had a single thread of execution: main thread.
- Program exits when the main thread exits.

► Multithreaded:

- Program is organized as multiple and concurrent threads of execution.
- The main thread *spawns* multiple threads.
- The threads **may** communicate with one another.
- Advantages:
 - Improves performance
 - Improves responsiveness
 - Improves utilization
 - less overhead compared to multiple processes

► Multithreaded programming

- Even in C, multithread programming may be accomplished in several ways
 - Pthreads: POSIX C library.
 - OpenMP
 - CUDA (GPU)
 - OpenCL (GPU/CPU)

► Example

```
int balance =500;
void deposit ( int sum ) {
    int currbalance=balance ; /* read balance */
    ...
    currbalance+=sum ;
    balance=currbalance ; /* write balance */
}
void withdraw ( int sum ) {
    int currbalance=balance ; /* read balance */
    if ( currbalance >0)
        currbalance-=sum ;
    balance=currbalance ; /* write balance */
}
.. deposit (100); /* thread 1*/
.. withdraw (50); /* thread 2*/
.. withdraw (100); /* thread 3*/
...
```

- minimize use of global/static memory
- Scenario: T1(read),T2(read,write),T1(write) ,balance=600
- Scenario: T2(read),T1(read,write),T2(write) ,balance=450

▶ Pthread

- ▶ Unlike Java, multithreading is not supported by the language standard.
- ▶ POSIX (Portable Operating System Interface, specified by IEEE)
- ▶ Threads (or Pthreads) is a POSIX standard for threads.
 - ▶ Included in the gcc compiler implementation

► Creating threads

```
int pthread_create(pthread_t * thread,  
                  const pthread_attr_t * attr,  
                  void *(*start_routine)(void*), void * arg);
```

- creates a new thread with the attributes specified by attr.
- Default attributes are used if attr is NULL.
 - E.g. set private stack size
- calls function start_routine(arg) on a separate thread of execution.
- returns zero on success, non-zero on error.

```
void pthread_exit(void *value_ptr);
```

- called implicitly when thread function exits.
-

► Example

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid)
{
    long tid;
    tid = (long)threadid;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t;
    for(t=0;t<NUM_THREADS;t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    /* Last thing that main() should do */
    pthread_exit(NULL);
}
```

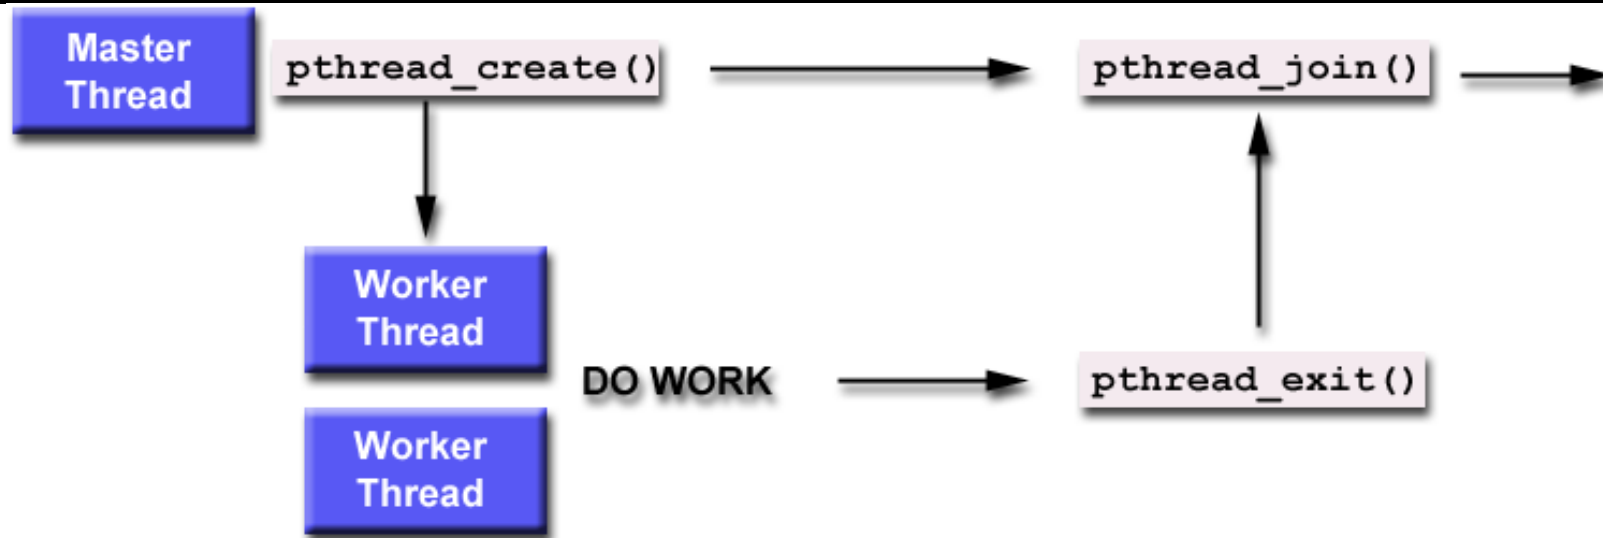
► Possible output

```
In main: creating thread 0
In main: creating thread 1
Hello World! It's me, thread #0!
Hello World! It's me, thread #1!
In main: creating thread 2
In main: creating thread 3
Hello World! It's me, thread #2!
Hello World! It's me, thread #3!
In main: creating thread 4
Hello World! It's me, thread #4!
```

```
In main: creating thread 0
Hello World! It's me, thread #0!
In main: creating thread 1
Hello World! It's me, thread #1!
In main: creating thread 2
Hello World! It's me, thread #2!
In main: creating thread 3
Hello World! It's me, thread #3!
In main: creating thread 4
Hello World! It's me, thread #4!
```

► Synchronization

Joining



int pthread_join(pthread_t thread, **void** **value_ptr);

- pthread_join() blocks the calling thread until the specified thread terminates.
- If value_ptr is not null, it will contain the return status of the ending thread.
- Other ways to synchronize: mutex, condition variables

```
▶ #define _OPEN_THREADS
▶ #include <pthread.h>
▶ #include <stdlib.h>
▶ #include <stdio.h>
▶
▶ void *thread(void *arg) {
▶     char *ret;
▶     printf("thread() entered with argument '%s'\n", arg);
▶     if ((ret = (char*) malloc(20)) == NULL) {
▶         perror("malloc() error");
▶         exit(2);
▶     }
▶     strcpy(ret, "This is a test");
▶     pthread_exit(ret);
▶ }
▶
▶ main() {
▶     pthread_t thid;
▶     void *ret;
▶
▶     if (pthread_create(&thid, NULL, thread, "thread 1") != 0) {
▶         perror("pthread_create() error");
▶         exit(1);
▶     }
▶
▶     if (pthread_join(thid, &ret) != 0) {
▶         perror("pthread_create() error");
▶         exit(3);
▶     }
▶
▶     printf("thread exited with '%s'\n", ret);
▶ }
```

```
pthread_exit((void*) t);
} Example
```

```
int main (int argc, char *argv[])
{
    pthread_t thread[NUM_THREADS];
    pthread_attr_t attr;
    int rc;
    long t;
    void *status;

    /* Initialize and set thread detached attribute */
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

    for(t=0; t<NUM_THREADS; t++) {
        printf("Main: creating thread %ld\n", t);
        rc = pthread_create(&thread[t], &attr, BusyWork, (void *)t);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    /* Free attribute and wait for the other threads */
    pthread_attr_destroy(&attr);
    for(t=0; t<NUM_THREADS; t++) {
        rc = pthread_join(thread[t], &status);
        if (rc) {
            printf("ERROR; return code from pthread_join() is %d\n", rc);
            exit(-1);
        }
        printf("Main: completed join with thread %ld having a status of %ld\n", t, (long)status);
    }

    printf("Main: program completed. Exiting.\n");
    pthread_exit(NULL);
}
```

► Mutex

- Mutex (mutual exclusion) acts as a "lock" protecting access to the shared resource.
- Only one thread can "own" the mutex at a time. Threads must take turns to lock the mutex.

```
int pthread_mutex_destroy(pthread_mutex_t *mutex);  
int pthread_mutex_init(pthread_mutex_t * mutex,  
                        const pthread_mutexattr_t * attr);  
pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;
```

- `pthread_mutex_init()` initializes a mutex. If attributes are NULL, default attributes are used.
- The macro `PTHREAD_MUTEX_INITIALIZER` can be used to initialize static mutexes.
- `pthread_mutex_destroy()` destroys the mutex.
- Both function return return 0 on success, non zero on error.

► Mutex

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- `pthread_mutex_lock()` locks the given mutex. If the mutex is locked, the function is blocked until it becomes available.
- `pthread_mutex_trylock()` is the non-blocking version. If the mutex is currently locked the call will return immediately.
- `pthread_mutex_unlock()` unlocks the mutex.

► Example revisted

```
int balance=500;
void deposit(int sum){
    int currbalance=balance; /*read balance*/
    ...
    currbalance+=sum;
    balance=currbalance; /*write balance*/
}

void withdraw(int sum){
    int currbalance=balance; /*read balance*/
    if (currbalance>0)
        currbalance-=sum;
    balance=currbalance; /*write balance*/
}

..
deposit(100); /*thread 1*/
..
withdraw(50); /*thread 2*/
..
withdraw(100); /*thread 3*/
...
```

- Scenario: T1(read),T2(read,write),T1(write),balance=600
- Scenario: T2(read),T1(read,write),T2(write),balance=450

► Using Mutex

```
int balance=500;
pthread_mutex_t mutexbalance=PTHREAD_MUTEX_INITIALIZER;

void deposit(int sum){
    pthread_mutex_lock(&mutexbalance);
    {
        int currbalance=balance; /* read balance */
        ...
        currbalance+=sum;
        balance=currbalance; /* write balance */
    }
    pthread_mutex_unlock(&mutexbalance);
}

void withdraw(int sum){
    pthread_mutex_lock(&mutexbalance);
    {
        int currbalance=balance; /* read balance */
        if (currbalance>0)
            currbalance-=sum;
        balance=currbalance; /* write balance */
    }
    pthread_mutex_unlock(&mutexbalance);
}

.. deposit(100); /* thread 1 */
.. withdraw(50); /* thread 2 */
.. withdraw(100); /* thread 3 */
```

- Scenario: T1(read,write),T2(read,write),balance=450
- Scenario: T2(read),T1(read,write),T2(write),balance=450

► Condition variables

- Sometimes locking or unlocking is based on a run-time condition.
 - Without condition variables, program would have to poll the variable/condition continuously.

Consumer:

- (a) lock mutex on global item variable
 - (b) wait for (item>0) signal from producer (mutex unlocked automatically).
 - (c) wake up when signaled (mutex locked again)
- automatically), unlock mutex and proceed.

Producer:

- (1) produce something
- (2) Lock global item variable, update item
- (3) signal waiting (threads)
- (4) unlock mutex

► Condition variables

```
int pthread_cond_destroy(pthread_cond_t *cond);  
int pthread_cond_init(pthread_cond_t * cond, const pthread_condattr_t * attr );  
pthread_cond_t cond = PTHREAD_COND_INITIALIZER;
```

- `pthread_cond_init()` initialized the condition variable. If `attr` is `NULL`, default attributes are set.
- `pthread_cond_destroy()` will destroy (uninitialize) the condition variable.
- destroying a condition variable upon which other threads are currently blocked results in undefined behavior.
- macro `PTHREAD_COND_INITIALIZER` can be used to initialize condition variables. No error checks are performed.
- Both function return 0 on success and non-zero otherwise.

► Condition variables

`int pthread_cond_wait(pthread_cond_t *cond, pthread_mutex_t *mutex);`

- blocks on a condition variable.
- must be called with the mutex already locked otherwise behavior undefined.
- automatically releases mutex
- upon successful return, the mutex will be automatically locked again.

`int pthread_cond_broadcast(pthread_cond_t *cond);`

`int pthread_cond_signal(pthread_cond_t *cond);`

- unblocks threads waiting on a condition variable.
- `pthread_cond_broadcast()` unlocks all threads that are waiting.
- `pthread_cond_signal()` unlocks one of the threads that are waiting.
- both return 0 on success, non zero otherwise.

► Example

```
#include <pthread.h>
pthread_cond_t cond_recv=PTHREAD_COND_INITIALIZER ;
pthread_cond_t cond_send=PTHREAD_COND_INITIALIZER ;
pthread_mutex_t cond_mutex=PTHREAD_MUTEX_INITIALIZER;
pthread_mutex_t count_mutex=PTHREAD_MUTEX_INITIALIZER;
```

```
int full=0;
int count=0;
```

```
void* produce ( void *)
{
    while(1)
    {
        pthread_mutex_lock(&cond_mutex);
        while (full)
        {
            pthread_cond_wait(&cond_recv,
                             &cond_mutex );
        }
        pthread_mutex_unlock(&cond_mutex);
        pthread_mutex_lock(&count_mutex);
        count++; full=1;
        printf("produced(%d):%d\n",
            pthread_self(), count);
        pthread_cond_broadcast(&cond_send);
        pthread_mutex_unlock(&count_mutex);
        if (count >=10) break;
    }
}
```

```
void* consume ( void *)
{
    while(1)
    {
        pthread_mutex_lock(&cond_mutex);
        while (!full)
        {
            pthread_cond_wait(&cond_send,
                             &cond_mutex );
        }
        pthread_mutex_unlock(&cond_mutex);
        pthread_mutex_lock(&count_mutex);
        full=0;
        printf("consumed(%ld):%d\n",
            pthread_self(), count);
        pthread_cond_broadcast(&cond_recv);
        pthread_mutex_unlock(&count_mutex);
        if (count >=10) break;
    }
}
```

► Example

```
int main ()
{
    pthread_t cons_thread, prod_thread;
    pthread_create(&prod_thread, NULL, produce, NULL);
    pthread_create(&cons_thread, NULL, consume, NULL);

    pthread_join(cons_thread, NULL);
    pthread_join(prod_thread, NULL);
    return 0;
}
```

Output:

```
produced(3077516144):1
consumed(3069123440):1
produced(3077516144):2
consumed(3069123440):2
produced(3077516144):3
consumed(3069123440):3
produced(3077516144):4
consumed(3069123440):4
produced(3077516144):5
consumed(3069123440):5
produced(3077516144):6
consumed(3069123440):6
produced(3077516144):7
consumed(3069123440):7
```

▶ Race conditions

- ▶ Race conditions occur when multiple threads share a variable, without proper synchronization
- ▶ Synchronization uses special variables, like a mutex, to ensure order of execution is correct
- ▶ Example: thread T1 needs to do something before thread T2
 - ▶ condition variable forces thread T2 to wait for thread T1
 - ▶ producer-consumer model program
- ▶ Example: two threads both need to access a variable and modify it based on its value
 - ▶ surround access and modification with a mutex
 - ▶ mutex groups operations together to make them atomic – treated as one unit

► Race conditions

Examples

- Consider the following program `race.c`:
 - What is the value of `cnt`?

```
unsigned int cnt = 0;
```

```
void *count ( void *arg ) { /* thread body */  
    int i;  
    for ( i = 0 ; i < 1000000000 ; i ++ )  
        cnt ++;  
    return NULL;  
}
```

```
int main ( void ) {  
    pthread_t tids [4];  
    int i;  
    for ( i = 0 ; i < 4 ; i ++ )  
        pthread_create ( &tids [i], NULL, count , NULL );  
    for ( i = 0 ; i < 4 ; i ++ )  
        pthread_join ( tids [i], NULL );  
    printf ( "cnt=%u\n" , cnt );  
    return 0;  
}
```

► Race conditions

Example

► Ideally, should increment `cnt` 4×100000000 times, so `cnt` = 400000000.

► However, running our code gives:

► `./race.o cnt=137131900`

► `./race.o cnt=163688698`

► `./race.o cnt=163409296`

► `./race.o cnt=170865738`

► `./race.o cnt=169695163`

► So, what happened?

▶ Race conditions

On assembly level

- ▶ C not designed for multithreading
- ▶ No notion of atomic operations in C
- ▶ Increment `cnt++`; maps to three assembly operations:
 - ▶ load `cnt` into a register
 - ▶ increment value in register
 - ▶ save new register value as new `cnt`
- ▶ So what happens if thread interrupted in the middle?
- ▶ Race condition!

► Race conditions

Fixed example

```
pthread_mutex_t mutex ;
unsigned i n t cnt = 0 ;

void *count ( void *arg ) { /* thread body */
    i n t i ;
    for ( i = 0 ; i < 1000000000 ; i ++ ) {
        pthread_mutex_lock ( & mutex ) ;
        cnt ++ ;
        pthread_mutex_unlock ( & mutex ) ;
    }
    return NULL ;
}


i n t main ( void ) {
    pthread_t tids [ 4 ] ;
    i n t i ;
    pthread_mutex_init ( & mutex , NULL ) ;
    for ( i = 0 ; i < 4 ; i ++ )
        pthread_create ( & tids [ i ] , NULL , count , NULL ) ;
    for ( i = 0 ; i < 4 ; i ++ )
        pthread_join ( tids [ i ] , NULL ) ; pthread_mutex_destroy ( & mutex ) ;
    printf ( "cnt=%u\n" , cnt ) ;
    return 0 ;
}
```

► Race conditions

- Note that new code functions correctly, but is much slower
- C statements are not atomic – threads may be interrupted at assembly level, in the middle of a C statement
- Atomic operations like mutex locking must be specified as atomic using special assembly instructions
- Ensure that all statements accessing/modifying shared variables are synchronized

► Example

- We consider a botanical garden that can be entered through two entrances (West and East)
- For the sake of simplification, visitors cannot leave the garden again.
- At each entrance, visitors must pass a turnstile
- Each turnstile has a counting device that counts up a common counter.



- Develop a C Program including Pthreads!